

Introduction:

You can use this ipython notebook as a template for the rest of the homework.

0. Basic Useful Setups & Read MNIST using Keras

```
In [1]: ### Basic setups
import cv2
import numpy as np
import matplotlib.pyplot as plt
import sys
sys.path.append('./models/')
# Enable automatic reload of libraries
%load_ext autoreload
# All modules are reloaded before every comment
%autoreload 2
import keras
from utils import load_mnist
(x_train, y_train), (x_test, y_test) = load_mnist()
from keras.utils.vis_utils import plot_model # can be used to plot the model into a png file.
```

```
/media/qiaojun/disk/Qiaojun/virtualenv/tensorflow/lib/python3.5/site-packages/h5py/__init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
```

```
from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

```
In [2]: #x_train_1 = np.reshape(x_train, (-1, 28*28))
y_train_1 = np.where(y_train==1)[1]
#x_test_1 = np.reshape(x_test, (-1, 28*28))
y_test_1 = np.where(y_test==1)[1]
x_train = np.squeeze(np.stack((x_train,)*3, axis = 3), axis=-1)
x_test = np.squeeze(np.stack((x_test,)*3, axis = 3), axis=-1)
```

```
In [3]: x_train_VGG = np.zeros((60000, 224, 224, 3), dtype = np.float32)
for index in range(60000):
    x_train_VGG[index, :, :, :] = cv2.resize(x_train[index, :, :, :], (224, 224))

x_test_VGG = np.zeros((10000, 224, 224, 3), dtype = np.float32)
for index in range(10000):
    x_test_VGG[index, :, :, :] = cv2.resize(x_test[index, :, :, :], (224, 224))
```

1. Convolutional Neural Network

1.1. LetNet Model

```
In [3]: from LeNet import LeNet
NUM_CLASSES = 10
model = LeNet(num_classes = NUM_CLASSES)
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])
# You can save the model by specifying a path
# modelPath = _____
# model.save(modelPath)

# You can also print out the model by typing
model.summary()

# You can plot the model into a png file and use it in your report.
pngModelPath = "images/LeNet.png"
# plot(model, to_file=pngModelPath, show_shapes=True)
plot_model(model, to_file=pngModelPath, show_shapes=True)
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 28, 28, 1)	0
block1_conv1 (Conv2D)	(None, 28, 28, 20)	520
block1_pool (MaxPooling2D)	(None, 14, 14, 20)	0
block2_conv2 (Conv2D)	(None, 14, 14, 50)	25050
block2_pool (MaxPooling2D)	(None, 7, 7, 50)	0
flatten (Flatten)	(None, 2450)	0
fc1 (Dense)	(None, 500)	1225500
predications (Dense)	(None, 10)	5010
Total params: 1,256,080		
Trainable params: 1,256,080		
Non-trainable params: 0		

```
In [4]: # train model
BATCH_SIZE = 128
MAX_EPOCH = 12

model.fit(x_train, y_train,
          batch_size=BATCH_SIZE,
          epochs=MAX_EPOCH,
          verbose=1,
          validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 9s 149us/step - loss: 0.1833 - acc: 0.9444 - val_loss: 0.0517 - val_acc: 0.9832

Epoch 2/12

60000/60000 [=====] - 3s 43us/step - loss: 0.0462 - acc: 0.9854 - val_loss: 0.0464 - val_acc: 0.9830

Epoch 3/12

60000/60000 [=====] - 3s 44us/step - loss: 0.0307 - acc: 0.9904 - val_loss: 0.0369 - val_acc: 0.9882

Epoch 4/12

60000/60000 [=====] - 3s 46us/step - loss: 0.0220 - acc: 0.9932 - val_loss: 0.0245 - val_acc: 0.9916

Epoch 5/12

60000/60000 [=====] - 3s 42us/step - loss: 0.0169 - acc: 0.9947 - val_loss: 0.0224 - val_acc: 0.9927

Epoch 6/12

60000/60000 [=====] - 3s 45us/step - loss: 0.0120 - acc: 0.9964 - val_loss: 0.0204 - val_acc: 0.9930

Epoch 7/12

60000/60000 [=====] - 3s 43us/step - loss: 0.0099 - acc: 0.9969 - val_loss: 0.0200 - val_acc: 0.9944

Epoch 8/12

60000/60000 [=====] - 3s 44us/step - loss: 0.0074 - acc: 0.9979 - val_loss: 0.0227 - val_acc: 0.9922

Epoch 9/12

60000/60000 [=====] - 3s 44us/step - loss: 0.0055 - acc: 0.9985 - val_loss: 0.0199 - val_acc: 0.9939

Epoch 10/12

60000/60000 [=====] - 2s 38us/step - loss: 0.0042 - acc: 0.9987 - val_loss: 0.0276 - val_acc: 0.9927

Epoch 11/12

60000/60000 [=====] - 2s 40us/step - loss: 0.0033 - acc: 0.9989 - val_loss: 0.0242 - val_acc: 0.9933

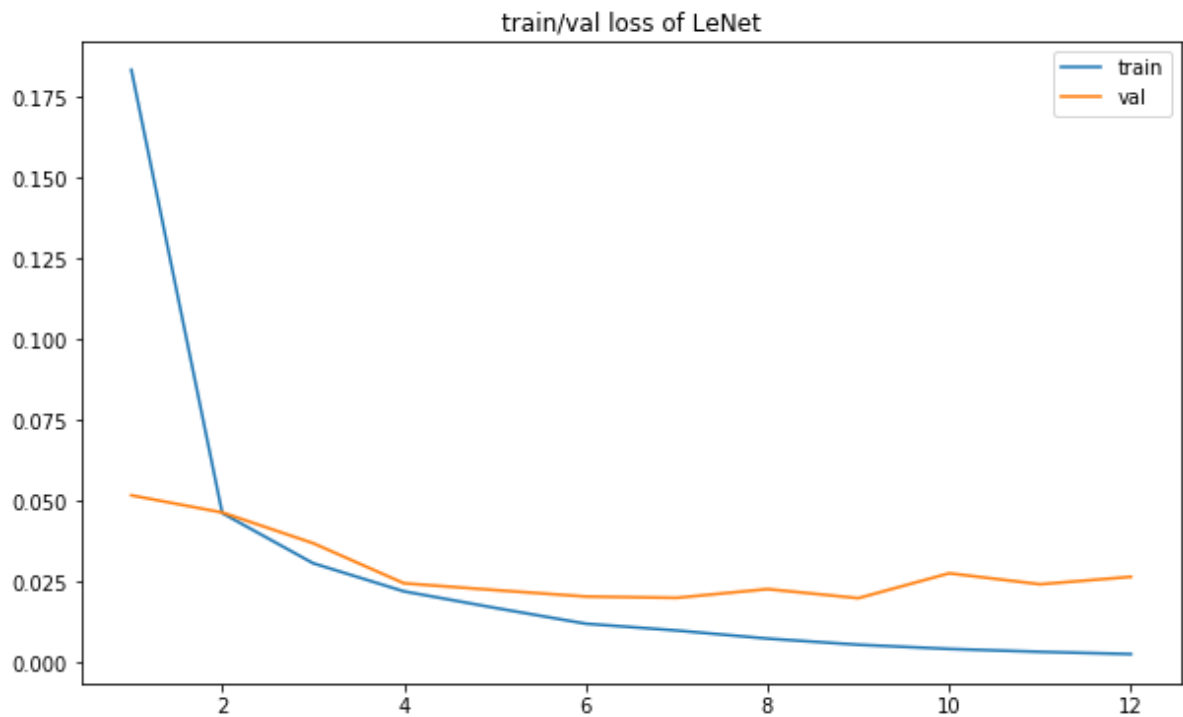
Epoch 12/12

60000/60000 [=====] - 3s 44us/step - loss: 0.0026 - acc: 0.9992 - val_loss: 0.0265 - val_acc: 0.9930

Test loss: 0.026506232323426637

Test accuracy: 0.993

```
In [10]: train_loss_LeNet = [0.1833,0.0462,0.0307,0.0220,0.0169,0.0120,0.0099,0.0074,0.0055,0.0042,0.0033,0.0026]
val_loss_LeNet = [0.0517,0.0464,0.0369,0.0245,0.0224,0.0204,0.0200,0.0227,0.0199,0.0276,0.0242,0.0265]
plt.figure(figsize=(10,6))
plt.plot(range(1,len(train_loss_LeNet)+1),train_loss_LeNet)
plt.plot(range(1,len(val_loss_LeNet)+1),val_loss_LeNet)
plt.legend(['train','val'])
plt.title('train/val loss of LeNet')
plt.show()
```



1.2. VGG16

```

In [4]: # train model
        BATCH_SIZE = 128
        MAX_EPOCH = 12

        from keras.applications.vgg16 import VGG16
        from keras.preprocessing import image
        from keras.applications.vgg16 import preprocess_input
        from keras.models import Model
        from keras.layers import Dense, GlobalAveragePooling2D
        import numpy as np

        # create the base pre-trained model
        base_model = VGG16(weights='imagenet', include_top=False)

        # add a global spatial average pooling layer
        x = base_model.output
        x = GlobalAveragePooling2D()(x)
        # let's add a fully-connected layer
        x = Dense(128, activation='relu')(x)
        # and a logistic layer -- let's say we have 10 classes
        predictions = Dense(10, activation='softmax')(x)

        # this is the model we will train
        model = Model(inputs=base_model.input, outputs=predictions)

        # first: train only the top layers (which were randomly initialized)
        # i.e. freeze all convolutional InceptionV3 layers
        for layer in base_model.layers:
            layer.trainable = False

        # compile the model (should be done *after* setting layers to non-trainable)
        # model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
        model.compile(loss=keras.losses.categorical_crossentropy,
                      optimizer=keras.optimizers.Adadelta(),
                      metrics=['accuracy'])
        # You can also print out the model by typing
        model.summary()

        # You can plot the model into a png file and use it in your report.
        #pngModelPath = "images/VGG16.png"
        # plot(model, to_file=pngModelPath, show_shapes=True)
        #plot_model(model, to_file=pngModelPath, show_shapes=True)

```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, None, None, 3)	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
global_average_pooling2d_1 ((None, 512)	0
dense_1 (Dense)	(None, 64)	32832
dense_2 (Dense)	(None, 10)	650
Total params: 14,748,170		
Trainable params: 33,482		
Non-trainable params: 14,714,688		

```
In [5]: model.fit(x_train_VGG, y_train,
                  batch_size=BATCH_SIZE,
                  epochs=MAX_EPOCH,
                  verbose=1,
                  validation_data=(x_test_VGG, y_test))
score = model.evaluate(x_test_VGG, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 667s 11ms/step - loss: 1.0711 - acc: 0.7268 - val_loss: 0.5223 - val_acc: 0.8770

Epoch 2/12

60000/60000 [=====] - 693s 12ms/step - loss: 0.3962 - acc: 0.9013 - val_loss: 0.2974 - val_acc: 0.9251

Epoch 3/12

60000/60000 [=====] - 865s 14ms/step - loss: 0.2664 - acc: 0.9315 - val_loss: 0.2245 - val_acc: 0.9406

Epoch 4/12

60000/60000 [=====] - 823s 14ms/step - loss: 0.2087 - acc: 0.9439 - val_loss: 0.1741 - val_acc: 0.9551

Epoch 5/12

60000/60000 [=====] - 931s 16ms/step - loss: 0.1760 - acc: 0.9512 - val_loss: 0.1817 - val_acc: 0.9447

Epoch 6/12

60000/60000 [=====] - 843s 14ms/step - loss: 0.1550 - acc: 0.9566 - val_loss: 0.1406 - val_acc: 0.9596

Epoch 7/12

60000/60000 [=====] - 940s 16ms/step - loss: 0.1398 - acc: 0.9609 - val_loss: 0.1260 - val_acc: 0.9648

Epoch 8/12

60000/60000 [=====] - 845s 14ms/step - loss: 0.1285 - acc: 0.9643 - val_loss: 0.1241 - val_acc: 0.9643

Epoch 9/12

60000/60000 [=====] - 946s 16ms/step - loss: 0.1191 - acc: 0.9657 - val_loss: 0.1148 - val_acc: 0.9666

Epoch 10/12

60000/60000 [=====] - 854s 14ms/step - loss: 0.1118 - acc: 0.9681 - val_loss: 0.1162 - val_acc: 0.9644

Epoch 11/12

60000/60000 [=====] - 821s 14ms/step - loss: 0.1063 - acc: 0.9695 - val_loss: 0.1046 - val_acc: 0.9675

Epoch 12/12

60000/60000 [=====] - 709s 12ms/step - loss: 0.1005 - acc: 0.9704 - val_loss: 0.1298 - val_acc: 0.9571

Test loss: 0.12977995154857636

Test accuracy: 0.9571

```
In [6]: for layer in model.layers[-6:]:
        layer.trainable = True
        from keras.optimizers import SGD
        model.compile(loss=keras.losses.categorical_crossentropy,
                      optimizer=SGD(lr=0.0001, momentum=0.9),
                      metrics=['accuracy'])
        model.fit(x_train_VGG, y_train,
                  batch_size=BATCH_SIZE,
                  epochs=MAX_EPOCH,
                  verbose=1,
                  validation_data=(x_test_VGG, y_test))
        score = model.evaluate(x_test_VGG, y_test, verbose=0)
        print('Test loss:', score[0])
        print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 757s 13ms/step - loss: 0.0787 - acc: 0.9759 - val_loss: 0.0656 - val_acc: 0.9785

Epoch 2/12

60000/60000 [=====] - 723s 12ms/step - loss: 0.0560 - acc: 0.9824 - val_loss: 0.0531 - val_acc: 0.9831

Epoch 3/12

60000/60000 [=====] - 785s 13ms/step - loss: 0.0472 - acc: 0.9851 - val_loss: 0.0472 - val_acc: 0.9851

Epoch 4/12

60000/60000 [=====] - 741s 12ms/step - loss: 0.0416 - acc: 0.9875 - val_loss: 0.0434 - val_acc: 0.9860

Epoch 5/12

60000/60000 [=====] - 801s 13ms/step - loss: 0.0371 - acc: 0.9883 - val_loss: 0.0416 - val_acc: 0.9871

Epoch 6/12

60000/60000 [=====] - 740s 12ms/step - loss: 0.0335 - acc: 0.9901 - val_loss: 0.0409 - val_acc: 0.9878

Epoch 7/12

60000/60000 [=====] - 801s 13ms/step - loss: 0.0309 - acc: 0.9909 - val_loss: 0.0398 - val_acc: 0.9876

Epoch 8/12

60000/60000 [=====] - 744s 12ms/step - loss: 0.0287 - acc: 0.9916 - val_loss: 0.0359 - val_acc: 0.9885

Epoch 9/12

60000/60000 [=====] - 788s 13ms/step - loss: 0.0264 - acc: 0.9925 - val_loss: 0.0355 - val_acc: 0.9887

Epoch 10/12

60000/60000 [=====] - 736s 12ms/step - loss: 0.0249 - acc: 0.9927 - val_loss: 0.0343 - val_acc: 0.9900

Epoch 11/12

60000/60000 [=====] - 803s 13ms/step - loss: 0.0238 - acc: 0.9931 - val_loss: 0.0353 - val_acc: 0.9892


Epoch 12/12

60000/60000 [=====] - 744s 12ms/step - loss: 0.0221 - acc: 0.9937 - val_loss: 0.0322 - val_acc: 0.9902

Test loss: 0.03215658742006635

Test accuracy: 0.9902


```
In [11]: train_loss_VGG = [1.0711,0.3962,0.2664,0.2087,0.1760,0.1550,0.1398,0.128
5,0.1191,0.1118,0.1063,0.1005,0.0787,0.0560,0.0472,0.0416,0.0371,0.0335,
0.0309,0.0287,0.0264,0.0249,0.0238,0.0221]
val_loss_VGG = [0.5223,0.2974,0.2245,0.1741,0.1817,0.1406,0.1260,0.1241,
0.1148,0.1162,0.1046,0.1298,0.0656,0.0531,0.0472,0.0434,0.0416,0.0409,0.
0398,0.0359,0.0355,0.0343,0.0353,0.0322]
plt.figure(figsize=(10,6))
plt.plot(range(1,len(train_loss_VGG)+1),train_loss_VGG)
plt.plot(range(1,len(val_loss_VGG)+1),val_loss_VGG)
plt.legend(['train','val'])
plt.title('train/val loss of VGG16')
plt.show()
```



2. 1-Nearest Neighbor

Using the raw pixel values (28×28) as feature vector, use 1-Nearest Neighbor to train a model and report its performance on the test set and the confusion matrix. Also report other details of the experiments like the distance functions used, prototype selection (if you have used), etc.

Use euclidean_distance (l2)

```
In [63]: from sklearn import neighbors

# the number of neighbors
n_neighbors = 1

# n_neighbors=5, weights='uniform', algorithm='auto', p=2, metric='minkowski', n_jobs=1
# weights: 'uniform', 'distance'
# p: distance norm
clf = neighbors.KNeighborsClassifier(n_neighbors, weights='uniform', n_jobs=-1)
clf.fit(x_train_1, y_train_1)

y_predict_1nn_1 = clf.predict(x_test_1)
acc_1nn_1 = sum(y_predict_1nn_1 == y_test_1)/len(y_test_1)
print("Test accuracy of 1-Nearest Neighbor: %.3f"%acc_1nn_1)
confusion_matrix(y_test_1, y_predict_1nn_1, labels=range(10))
```

Test accuracy of 1-Nearest Neighbor: 0.969

```
Out[63]: array([[ 973,    1,    1,    0,    0,    1,    3,    1,    0,    0],
 [    0, 1129,    3,    0,    1,    1,    1,    0,    0,    0],
 [    7,    6,  992,    5,    1,    0,    2,   16,    3,    0],
 [    0,    1,    2,  970,    1,   19,    0,    7,    7,    3],
 [    0,    7,    0,    0,  944,    0,    3,    5,    1,   22],
 [    1,    1,    0,   12,    2,  860,    5,    1,    6,    4],
 [    4,    2,    0,    0,    3,    5,  944,    0,    0,    0],
 [    0,   14,    6,    2,    4,    0,    0,  992,    0,   10],
 [    6,    1,    3,   14,    5,   13,    3,    4,  920,    5],
 [    2,    5,    1,    6,   10,    5,    1,   11,    1,  967]])
```

Use manhattan_distance (l1)

```
In [65]: from sklearn import neighbors
n_neighbors = 1
clf = neighbors.KNeighborsClassifier(n_neighbors, p=1, n_jobs=-1)
clf.fit(x_train_1, y_train_1)
y_predict_1nn_2 = clf.predict(x_test_1)
acc_1nn_2 = sum(y_predict_1nn_2 == y_test_1)/len(y_test_1)
print("Test accuracy of 1-Nearest Neighbor: %.3f"%acc_1nn_2)
confusion_matrix(y_test_1, y_predict_1nn_2, labels=range(10))
```

Test accuracy of 1-Nearest Neighbor: 0.963

```
Out[65]: array([[ 973,    2,    1,    0,    0,    1,    2,    1,    0,    0],
 [    0, 1129,    3,    0,    1,    1,    1,    0,    0,    0],
 [    9,    8,  987,    6,    1,    0,    2,   17,    2,    0],
 [    0,    2,    4,  965,    1,   21,    0,    9,    4,    4],
 [    1,    9,    0,    0,  937,    0,    3,    4,    1,   27],
 [    2,    1,    0,   17,    2,  848,    9,    1,    5,    7],
 [    5,    2,    1,    0,    2,    5,  943,    0,    0,    0],
 [    0,   20,    4,    2,    4,    0,    0,  989,    0,    9],
 [    9,    5,    6,   21,    4,   18,    3,    4,  894,   10],
 [    1,    5,    1,    7,   13,    5,    1,    9,    1,  966]])
```

3. Support Vector Machines

Using the same feature vector, use SVM to train a model and report its performance on the test set and the confusion matrix. Report other relevant details of the experiments.

```
In [ ]: # C : float, optional (default=1.0)
# Penalty parameter C of the error term.
# kernel : string, optional (default='rbf')
# Specifies the kernel type to be used in the algorithm.
# It must be one of 'linear', 'poly', 'rbf', 'sigmoid', 'precomputed'
# or a callable.
# If none is given, 'rbf' will be used.
# degree : int, optional (default=3)
# Degree of the polynomial kernel function ('poly'). Ignored by all other kernels.
# gamma : float, optional (default='auto')
# Kernel coefficient for 'rbf', 'poly' and 'sigmoid'. If gamma is 'auto' then 1/n_features will be used instead.

#SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
# decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
# max_iter=-1, probability=False, random_state=None, shrinking=True,
# tol=0.001, verbose=False)
```

Use Radial basis function(RBF) kernel

```
In [67]: from sklearn.svm import SVC
clf = SVC()
clf.fit(x_train_1, y_train_1)
y_predict_SVM_1 = clf.predict(x_test_1)
acc_SVM_1 = sum(y_predict_SVM_1 == y_test_1)/len(y_test_1)
print("Test accuracy of Support Vector Machines: %.3f"%acc_SVM_1)
confusion_matrix(y_test_1, y_predict_SVM_1, labels=range(10))
```

Test accuracy of Support Vector Machines: 0.945

```
Out[67]: array([[ 967,    0,    2,    0,    0,    5,    4,    1,    1,    0],
 [   0, 1121,    2,    2,    0,    1,    4,    1,    4,    0],
 [   9,    1,  961,    9,   10,    1,   13,    9,   17,    2],
 [   1,    1,   15,  951,    1,   15,    1,   10,   11,    4],
 [   1,    1,    7,    0,  938,    0,    7,    2,    2,   24],
 [   7,    4,    5,   30,    7,  810,   12,    2,   10,    5],
 [   9,    3,    4,    1,    5,    9,  926,    0,    1,    0],
 [   2,   13,   22,    5,    8,    1,    0,  956,    3,   18],
 [   4,    7,    7,   14,    8,   23,   10,    6,  892,    3],
 [   8,    7,    0,   12,   31,    6,    1,   13,    7,  924]])
```

Use linear kernel

```
In [70]: from sklearn.linear_model import SGDClassifier
clf = SGDClassifier(tol = 1e-3)
clf.fit(x_train_1, y_train_1)
y_predict_SVM_2 = clf.predict(x_test_1)
acc_SVM_2 = sum(y_predict_SVM_2 == y_test_1)/len(y_test_1)
print("Test accuracy of Support Vector Machines: %.3f"%acc_SVM_2)
confusion_matrix(y_test_1, y_predict_SVM_2, labels=range(10))
```

Test accuracy of Support Vector Machines: 0.918

```
Out[70]: array([[ 951,    0,    1,    2,    2,   11,    5,    6,    1,    1],
 [    0, 1108,    4,    0,    0,    2,    4,    2,   15,    0],
 [    4,    5,   939,   10,   12,    5,    9,   12,   33,    3],
 [    5,    0,   30,  893,    3,   29,    3,   13,   25,    9],
 [    1,    5,    5,    0,  931,    0,    6,    2,    8,   24],
 [    8,    4,    3,   28,   20,  770,   17,    4,   30,    8],
 [    5,    3,    7,    3,   10,   13,  907,    3,    7,    0],
 [    1,    7,   28,    3,   10,    2,    1,  947,    2,   27],
 [    6,    9,    8,   13,   14,   32,   10,   11,  859,   12],
 [    5,   11,    3,    8,   50,   12,    0,   34,   12,  874]])
```

4. Spatial Pyramid Matching

```
In [111]: %cd Image-recognition/  
          %run example.py  
          %cd ../
```

```

/media/qiaojun/disk/Qiaojun/Courses/COGS260/assignment/assignment2/hw2/
Image-recognition
Training data load from None
Test data load from None
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Codebook Size: 100
Pyramid level: 1
Dense SIFT feature extraction
Building the codebook, it will take some time
Spatial Pyramid Matching encoding
Tuning hyper-parameters

```

Best parameters set found on development set:

```

SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=100.0, kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)

```

Grid scores on development set:

```

0.112 (+/-0.000) for {'C': 0.001, 'gamma': 0.001}
0.112 (+/-0.000) for {'C': 0.001, 'gamma': 0.01}
0.112 (+/-0.000) for {'C': 0.001, 'gamma': 0.1}
0.112 (+/-0.000) for {'C': 0.001, 'gamma': 1.0}
0.433 (+/-0.004) for {'C': 0.001, 'gamma': 10.0}
0.612 (+/-0.013) for {'C': 0.001, 'gamma': 100.0}
0.112 (+/-0.000) for {'C': 0.01, 'gamma': 0.001}
0.112 (+/-0.000) for {'C': 0.01, 'gamma': 0.01}
0.112 (+/-0.000) for {'C': 0.01, 'gamma': 0.1}
0.525 (+/-0.009) for {'C': 0.01, 'gamma': 1.0}
0.726 (+/-0.015) for {'C': 0.01, 'gamma': 10.0}
0.775 (+/-0.015) for {'C': 0.01, 'gamma': 100.0}
0.112 (+/-0.000) for {'C': 0.1, 'gamma': 0.001}
0.112 (+/-0.000) for {'C': 0.1, 'gamma': 0.01}
0.529 (+/-0.010) for {'C': 0.1, 'gamma': 0.1}
0.730 (+/-0.015) for {'C': 0.1, 'gamma': 1.0}
0.794 (+/-0.012) for {'C': 0.1, 'gamma': 10.0}
0.835 (+/-0.013) for {'C': 0.1, 'gamma': 100.0}
0.112 (+/-0.000) for {'C': 1.0, 'gamma': 0.001}
0.530 (+/-0.010) for {'C': 1.0, 'gamma': 0.01}
0.730 (+/-0.015) for {'C': 1.0, 'gamma': 0.1}
0.790 (+/-0.014) for {'C': 1.0, 'gamma': 1.0}
0.829 (+/-0.010) for {'C': 1.0, 'gamma': 10.0}
0.862 (+/-0.010) for {'C': 1.0, 'gamma': 100.0}
0.530 (+/-0.010) for {'C': 10.0, 'gamma': 0.001}
0.730 (+/-0.016) for {'C': 10.0, 'gamma': 0.01}
0.789 (+/-0.015) for {'C': 10.0, 'gamma': 0.1}
0.816 (+/-0.010) for {'C': 10.0, 'gamma': 1.0}
0.852 (+/-0.011) for {'C': 10.0, 'gamma': 10.0}
0.858 (+/-0.010) for {'C': 10.0, 'gamma': 100.0}
0.731 (+/-0.015) for {'C': 100.0, 'gamma': 0.001}
0.789 (+/-0.015) for {'C': 100.0, 'gamma': 0.01}
0.812 (+/-0.010) for {'C': 100.0, 'gamma': 0.1}
0.835 (+/-0.010) for {'C': 100.0, 'gamma': 1.0}

```

0.860 (+/-0.010) for {'C': 100.0, 'gamma': 10.0}
 0.844 (+/-0.010) for {'C': 100.0, 'gamma': 100.0}

Detailed classification report:

The model is trained on the full development set.
 The scores are computed on the full evaluation set.

	precision	recall	f1-score	support
0	0.89	0.94	0.92	980
1	0.98	0.99	0.98	1135
2	0.87	0.86	0.87	1032
3	0.79	0.80	0.80	1010
4	0.92	0.88	0.90	982
5	0.82	0.78	0.80	892
6	0.93	0.93	0.93	958
7	0.85	0.88	0.86	1028
8	0.80	0.75	0.77	974
9	0.81	0.84	0.83	1009
avg / total	0.87	0.87	0.87	10000

Test accuracy of SPM: 0.868

```
[[ 922   3   3   1   5   5  19   5  12   5]
 [   1 1119   2   3   0   1   4   1   4   0]
 [   5   2 888  79   6   5   2  33  10   2]
 [   0   0  74 813   2  30   0  61  25   5]
 [  10   0   4   0 866   9  18   2  13  60]
 [  11   3   6  52   7 699   8  10  58  38]
 [  36   4   3   0   9   6 887   0  13   0]
 [   2   1  27  38   5   5   0 907   9  34]
 [  29   4  12  41  13  56   8  30 730  51]
 [  16   4   0   8  25  37   4  23  40 852]]
```

/media/qiaojun/disk/Qiaojun/Courses/COGS260/assignment/assignment2/hw2

In [113]: acc_SPM = 0.868

5. Deep Belief Nets

```
In [115]: #Test accuracy of Deep Belief Nets: 0.985
# [[ 975   1   0   0   0   1   2   1   0   0]
# [   0 1129   2   1   0   1   1   0   1   0]
# [   6   0 1017   1   1   0   1   4   2   0]
# [   1   0   3 991   0   4   0   5   4   2]
# [   2   0   2   0 957   0   5   1   0 15]
# [   2   0   0   1   1 878   4   0   4   2]
# [   2   2   0   1   2   2 945   0   4   0]
# [   2   3   6   0   1   0   0 1009   1   6]
# [   2   1   1   2   0   1   0   1 962   4]
# [   2   3   1   2   3   2   0   3   3 990]]
```

Evaluation

```
In [7]: name_list = ['LeNet', 'VGG16', '1-NN', 'SVM', 'SPM', 'DBN']  
acc_list = [0.993, 0.990, 0.969, 0.945, 0.868, 0.985]  
  
plt.figure(figsize=(10,6))  
plt.bar(range(len(name_list)), acc_list, color='rgbmc', tick_label=name_  
list)  
plt.title('Evaluation of classification accuracy on MNIST')  
plt.show()
```

