


# COGS 260 Assignment3

## 1 Least Square Estimation

```
In [134]: # 1.1 Import packages and load data
import numpy as np
import matplotlib.pyplot as plt
X_and_Y = np.load('./ql-least-square.npy')
X = X_and_Y[:, 0] # Shape: (300,)
Y = X_and_Y[:, 1] # Shape: (300,)
```


```
In [52]: # 1.2 Plot the scatter graph of data
plt.scatter(X, Y)
plt.xlabel('X')
plt.ylabel('Y')
plt.show()
```



```
In [40]: # 1.3 Compute the least square line over the given data.
# Assume  $Y = w_0 + w_1 * X = (w_0, w_1) \cdot (1, X) = W \cdot X_1$ 
#  $X_1$  contains 1 and  $X$ .
X1 = np.matrix(np.hstack((np.ones((len(X), 1)), X.reshape(-1, 1))))
W = X1.T.dot(X1).I.dot(X1.T).dot(Y)
w0, w1 = np.array(W).reshape(-1)
print('Y = {:.2f} + {:.2f}*X'.format(w0, w1))

Y = -15.47 + 11.61*X
```


```
In [42]: # 1.4 Plot the scatter graph of data and estimated line.
X_line = np.linspace(0,10,300)
Y_line = w0 + w1 * X_line
plt.scatter(X, Y)
plt.plot(X_line, Y_line, color='orange')
plt.xlabel('X')
plt.ylabel('Y')
plt.savefig('imgs/1_1.png')
plt.show()
```



```
In [33]: # 1.5 Compute the least square parabola over the given data.
# Assume  $Y = w_0 + w_1X + w_2X^2 = (w_0, w_1, w_2) \cdot (1, X, X^2) = W \cdot X_2$ 
#  $X_2$  contains 1, X and  $X^2$ .
X2 = np.matrix(np.hstack((np.ones((len(X),1)),X.reshape(-1,1)
))),X.reshape(-1,1)**2)))
W = X2.T.dot(X2).I.dot(X2.T).dot(Y)
w0, w1, w2 = np.array(W).reshape(-1)
print('Y = {:.2f} + {:.2f}*X + {:.2f}*X^2'.format(w0, w1, w2))

Y = -1.71 + 3.02*X + 0.87*X^2
```


```
In [35]: # 1.6 Plot the scatter graph of data and estimated parabola
X_line = np.linspace(0,10,300)
Y_line = w0 + w1 * X_line + w2 * (X_line**2)
plt.scatter(X, Y)
plt.plot(X_line, Y_line, color='orange')
plt.xlabel('X')
plt.ylabel('Y')
plt.savefig('imgs/1_2.png')
plt.show()
```



## 2 Parabola Estimation

```
In [109]: # 2.1 Import packages and load data
import numpy as np
import matplotlib.pyplot as plt
X_and_Y = np.load('./q2-parabola.npy')
X = X_and_Y[:, 0] # Shape: (300,)
Y = X_and_Y[:, 1] # Shape: (300,)
```

```
In [110]: # 2.2 Plot the scatter graph of data
plt.scatter(X, Y)
plt.xlabel('X')
plt.ylabel('Y')
plt.show()
```




We are comparing the robustness of different loss definition so we can see a bunch of outliers at the upleft corner.

```
In [111]: # 2.3 Compute the least square(L2 norm) parabola over the given data.
# Assume  $Y = w_0 + w_1 * X + w_2 * X^2 = (w_0, w_1, w_2) \cdot (1, X, X^2) = W \cdot X_2$ 
#  $X_2$  contains 1, X and  $X^2$ .
X1 = np.matrix(np.hstack((np.ones((len(X),1)),X.reshape(-1,1)
))),X.reshape(-1,1)**2)))
W = X1.T.dot(X1).I.dot(X1.T).dot(Y)
w0, w1, w2 = np.array(W).reshape(-1)
print('Y = {:.2f} + {:.2f}*X + {:.2f}*X^2'.format(w0, w1, w2))

Y = 51.07 + -16.06*X + 2.36*X^2
```

```
In [112]: # 2.4 Plot the scatter graph of data and estimated parabola
X_line = np.linspace(0,10,300)
Y_line = w0 + w1 * X_line + w2 * (X_line**2)
plt.scatter(X, Y)
plt.plot(X_line, Y_line, color='orange')
plt.xlabel('X')
plt.ylabel('Y')
plt.savefig('imgs/2_1.png')
plt.show()
```




```
In [124]: # 2.5 Compute the L1 norm loss parabola over the given data.
# Assume  $Y = w_0 + w_1 * X + w_2 * X^2 = (w_0, w_1, w_2) \cdot (1, X, X^2) = W \cdot X_2$ 
#  $X_2$  contains 1, X and  $X^2$ .
X2 = np.matrix(np.hstack((np.ones((len(X),1)),X.reshape(-1,1)
))),X.reshape(-1,1)**2)))
Y2 = np.expand_dims(Y,1)
# First randomly pick a W
W = np.random.normal(0,5,(3,1))
# Define the L1-norm loss
loss = np.sum(np.abs(X2.dot(W) - Y2))
# gradient descent parameters (refer to the hint)
num_point = len(Y)
max_iterations = 300000
learning_rate = 0.000001
threshold = 0.00001
iters = 0
while loss/num_point > threshold and iters < max_iterations:
    grad = np.sign(X2.dot(W) - Y2).T.dot(X2).T
    W = W - learning_rate*grad
    loss = np.sum(np.abs(X2.dot(W) - Y2))
    iters += 1
print("total iterations %d \t average loss %f"%(iters,loss/num_point))

w0, w1, w2 = np.array(W).reshape(-1)
print('Y = {:.2f} + {:.2f}*X + {:.2f}*X^2'.format(w0, w1, w2))

total iterations 300000          average loss 12.058130
Y = 2.48 + 0.77*X + 1.08*X^2
```

```
In [125]: # 2.6 Plot the scatter graph of data and estimated parabola
X_line = np.linspace(0,10,300)
Y_line = w0 + w1 * X_line + w2 * (X_line**2)
plt.scatter(X, Y)
plt.plot(X_line, Y_line, color='orange')
plt.xlabel('X')
plt.ylabel('Y')
plt.savefig('imgs/2_2.png')
plt.show()
```



### 3 Perceptron Learning

```
In [110]: # 3.1 Dataset
from sklearn import preprocessing
import numpy as np
import pandas as pd

# load iris_train.data and iris_test.data
pd_train = pd.read_csv("iris/iris_train.data", names=["sepal_length", "sepal_width", "petal_length", "petal_width", "label"])
pd_test = pd.read_csv("iris/iris_test.data", names=["sepal_length", "sepal_width", "petal_length", "petal_width", "label"])

# parse the features and labels as numpy arrays.
X_train = pd_train.as_matrix(columns=["sepal_length", "sepal_width", "petal_length", "petal_width"])
y_train = pd_train.as_matrix(columns=["label"]).ravel()
X_test = pd_test.as_matrix(columns=["sepal_length", "sepal_width", "petal_length", "petal_width"])
y_test = pd_test.as_matrix(columns=["label"]).ravel()

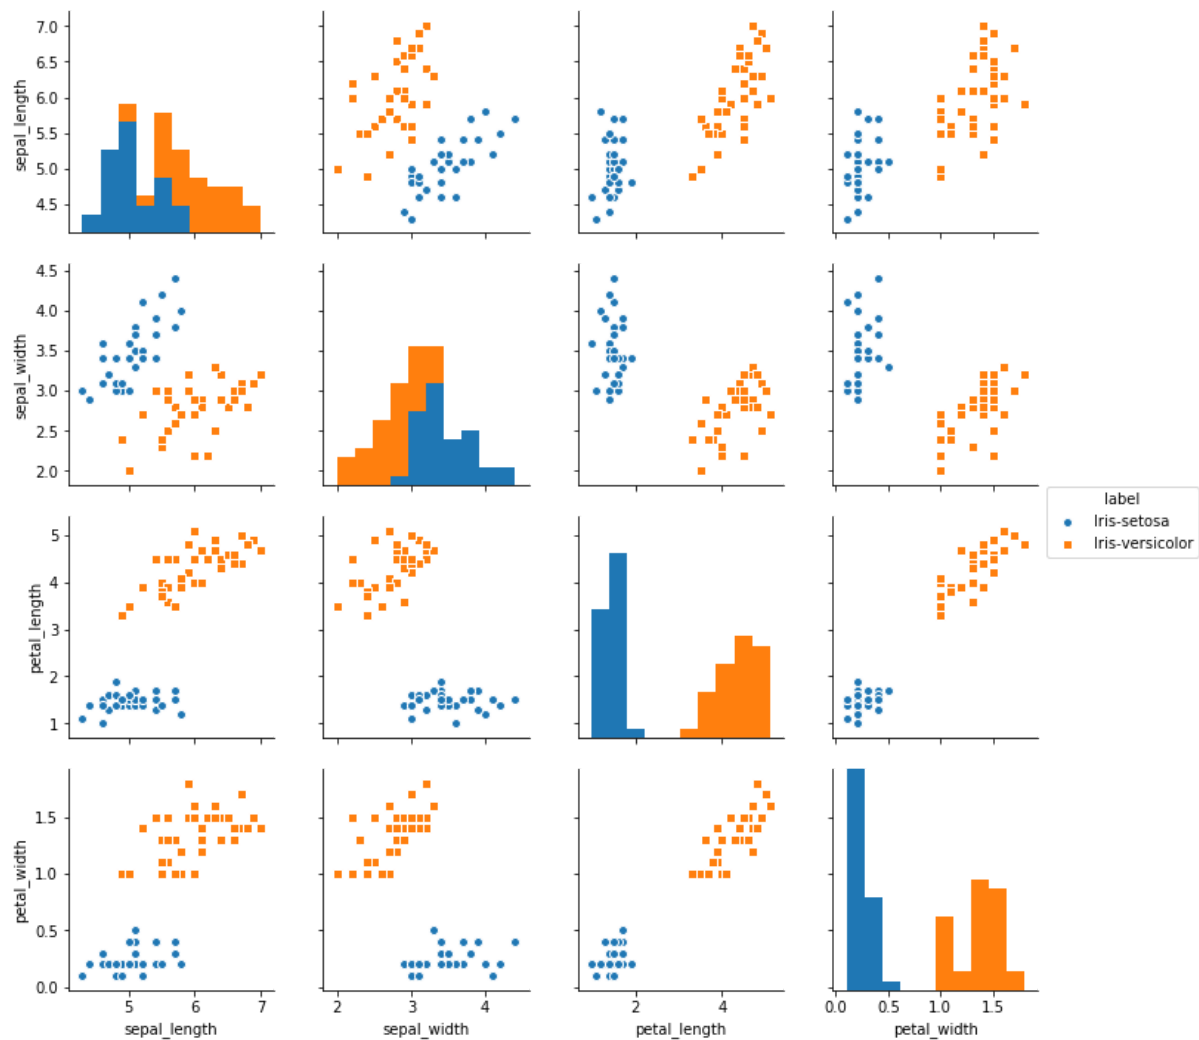
# Encode labels
le = preprocessing.LabelEncoder()
le.fit(y_train)
print(le.classes_)
y_train = le.transform(y_train)
y_test = le.transform(y_test)

['Iris-setosa' 'Iris-versicolor']
```

```
In [115]: type(pd_train)
```

```
Out[115]: pandas.core.frame.DataFrame
```

```
In [119]: import seaborn as sns  
g = sns.pairplot(pd_train,hue="label",markers=["o", "s"])
```






```
In [299]: # 3.2 Programming
W = np.random.normal(0,5,(4,1))
b = np.random.normal(0,5,1)
y_predict = (X_train.dot(W)+b>0).reshape(-1)
loss = np.sum(y_predict!=y_train)

num_point = len(y_train)
max_iterations = 300000
Lambda = 1
iters = 0
loss_list = []
while loss > 0 and iters < max_iterations:
    rand_int = np.random.randint(num_point)
    y_predict = X_train[rand_int].dot(W)+b>0
    if y_train[rand_int] == y_predict:
        continue
    else:
        W = W + Lambda*(y_train[rand_int]-y_predict)*np.expand_dims(X_train[rand_int],1)
        b = b + Lambda*(y_train[rand_int]-y_predict)
        y_predict = (X_train.dot(W)+b>0).reshape(-1)
        loss = np.sum(y_predict!=y_train)/num_point
        iters += 1
        loss_list.extend([loss])
        print("iter %2d \t loss %f"%(iters,loss))

plt.plot(range(1,1+iters),loss_list)
plt.xlabel('iters')
plt.ylabel('error rate')
plt.show()
```

iter	1	loss	1.000000
iter	2	loss	0.500000
iter	3	loss	0.500000
iter	4	loss	0.500000
iter	5	loss	0.500000
iter	6	loss	0.500000
iter	7	loss	0.500000
iter	8	loss	0.057143
iter	9	loss	0.500000
iter	10	loss	0.085714
iter	11	loss	0.500000
iter	12	loss	0.371429
iter	13	loss	0.500000
iter	14	loss	0.485714
iter	15	loss	0.014286
iter	16	loss	0.500000
iter	17	loss	0.000000



```
In [300]: # 3.3 Decision Boundary
w0, w1, w2, w3 = np.array(W).reshape(-1)
b = b[0]
print('Y = {:.2f}*sepal_length + {:.2f}*sepal_width + {:.2f}*petal_length + {:.2f}*petal_width + {:.2f}'.format(w0, w1, w2, w3, b))

Y = 2.14*sepal_length + -15.07*sepal_width + 7.66*petal_length + 6.91*petal_width + 2.32
```

```
In [301]: # 3.4 Test
# 1.accuracy
num_point = len(y_test)
y_predict = (X_test.dot(W)+b>0).reshape(-1)
accuracy = np.sum(y_predict == y_test) / num_point
print('accuracy:\t', accuracy)
# For the following, let y_test==1 be positive
true_positive = np.sum(np.logical_and(y_test==1,y_predict==1))
# 2.Precision
precision = true_positive/np.sum(y_predict)
print('precision:\t', precision)
# 3.Recall
recall = true_positive/np.sum(y_test)
print('recall:\t\t',recall)
# 4.F-value
F_value = 2*precision*recall / (precision + recall)
print('F-value:\t',F_value)
```

```
accuracy:      1.0
precision:     1.0
recall:        1.0
F-value:       1.0
```

```

In [287]: # 3.5 [bonus]Z-score
X_mean = np.mean(X_train,axis=0)
X_var = np.var(X_train,axis=0)

X_train_normalized = (X_train-X_mean)/X_var

W = np.random.normal(0,5,(4,1))
b = np.random.normal(0,5,1)
y_predict = (X_train_normalized.dot(W)+b>0).reshape(-1)
loss = np.sum(y_predict!=y_train)

num_point = len(y_train)
max_iterations = 300000
Lambda = 1
iters = 0
loss_list = []
while loss > 0 and iters < max_iterations:
    rand_int = np.random.randint(num_point)
    y_predict = X_train_normalized[rand_int].dot(W)+b>0
    if y_train[rand_int] == y_predict:
        continue
    else:
        W = W + Lambda*(y_train[rand_int]-y_predict)*np.expand_dims(X_train_normalized[rand_int],1)
        b = b + Lambda*(y_train[rand_int]-y_predict)
        y_predict = (X_train_normalized.dot(W)+b>0).reshape(-1)
        loss = np.sum(y_predict!=y_train)/num_point
        iters += 1
        loss_list.extend([loss])
        print("iter %2d \t loss %f"%(iters,loss))


plt.plot(range(1,1+iters),loss_list)
plt.xlabel('iters')
plt.ylabel('error rate')
plt.show()

w0, w1, w2, w3 = np.array(W).reshape(-1)
b = b[0]
print('Y = {:.2f}*sepal_length + {:.2f}*sepal_width + {:.2f}*petal_length + {:.2f}*petal_width + {:.2f}'.format(w0, w1, w2, w3, b))

num_point = len(y_test)
X_test_normalized = (X_test-X_mean)/X_var
y_predict = (X_test.dot(W)+b>0).reshape(-1)
accuracy = np.sum(y_predict == y_test) / num_point
print('accuracy:\t', accuracy)
# For the following, let y_test==1 be positive
true_positive = np.sum(np.logical_and(y_test==1,y_predict==1))
# 2.Precision
precision = true_positive/np.sum(y_predict)
print('precision:\t', precision)
# 3.Recall
recall = true_positive/np.sum(y_test)
print('recall:\t\t',recall)
# 4.F-value

```

```
F_value = 2*precision*recall / (precision + recall)
print('F-value: %f' % F_value)
iter  2      loss 0.285714
iter  3      loss 0.057143
iter  4      loss 0.000000
```



```
Y = 5.46*sepal_length + -3.56*sepal_width + 1.39*petal_length + 0.23*petal_width + 3.19
accuracy:      0.5
precision:     0.5
recall:        1.0
F-value:       0.6666666666666666
```

After normalizing the features, the weights on features have changed. Before `sepal_width` has the largest weight. Now `sepal_length` is the most important one.

Notice that with z-score features, the training process tends to converge faster. However, it is likely to overfit the training set since this dataset is pretty small so the distribution difference between training set and test set is not too small to ignore. That's why the test performance is not as good as using origin features

## 4 Feed Forward Neural Network

```
In [18]: # preparing the data
import numpy as np
import matplotlib.pyplot as plt
from utils import load_mnist
(x_train, y_train), (x_test, y_test) = load_mnist()
x_train_vector = np.reshape(x_train, (-1, 28*28))
x_test_vector = np.reshape(x_test, (-1, 28*28))
x_valid_vector = x_train_vector[50000:]
y_valid = y_train[50000:]
x_train_vector = x_train_vector[:50000]
y_train = y_train[:50000]
print('train num:%d valid num:%d test num:%d'%(len(y_train), len(y_valid),
len(y_test)))

def softmax(x):
    e_x = np.exp(x - np.max(x))
    return e_x / np.expand_dims(e_x.sum(axis=1), 1)

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
train num:50000 valid num:10000 test num:10000
```


```

In [22]: # 4.1 Train a feed forward network with 1 hidden layer
# Here we use Cross Entropy Loss with Softmax activation function
W = np.random.normal(0,0.1,(784,10))
num_point = len(y_train)
max_iterations = 1000
learning_rate = 1e-5
threshold = 0.0001
loss = num_point
iters = 0
loss_list = []
train_error_list = []
valid_error_list = []
while loss > threshold and iters < max_iterations:
    layer_output = x_train_vector.dot(W)
    y_predict = softmax(layer_output)
    loss = - np.sum(y_train * np.log(np.clip(y_predict,1e-15,np.inf)))/num_point
    grad = y_predict - y_train
    W_grad = x_train_vector.T.dot(grad)
    W -= learning_rate*W_grad
    iters += 1
    loss_list.extend([loss])
    if iters%20 == 0:
        y_max = np.argmax(y_predict,axis=1)
        train_error = 1-np.sum(y_train[range(len(y_train)),y_max])/len(y_train)
        train_error_list.extend([train_error])
        layer_output = x_valid_vector.dot(W)
        y_predict = softmax(layer_output)
        y_max = np.argmax(y_predict,axis=1)
        valid_error = 1-np.sum(y_valid[range(len(y_valid)),y_max])/len(y_valid)
        valid_error_list.extend([valid_error])
        print('iter %4d loss: %.3f train: %.3f valid: %.3f'%(iters,loss,train_error,valid_error))
plt.plot(range(1,1+iters),loss_list)
plt.xlabel('iters')
plt.ylabel('loss')
plt.show()

```


iter	20	loss:	0.655	train:	0.168	valid:	0.143
iter	40	loss:	0.517	train:	0.138	valid:	0.119
iter	60	loss:	0.463	train:	0.125	valid:	0.109
iter	80	loss:	0.432	train:	0.118	valid:	0.104
iter	100	loss:	0.411	train:	0.112	valid:	0.101
iter	120	loss:	0.396	train:	0.109	valid:	0.098
iter	140	loss:	0.385	train:	0.106	valid:	0.095
iter	160	loss:	0.375	train:	0.104	valid:	0.093
iter	180	loss:	0.368	train:	0.103	valid:	0.092
iter	200	loss:	0.361	train:	0.101	valid:	0.090
iter	220	loss:	0.356	train:	0.099	valid:	0.089
iter	240	loss:	0.351	train:	0.098	valid:	0.088
iter	260	loss:	0.347	train:	0.097	valid:	0.087
iter	280	loss:	0.343	train:	0.096	valid:	0.087
iter	300	loss:	0.340	train:	0.095	valid:	0.086
iter	320	loss:	0.337	train:	0.094	valid:	0.086
iter	340	loss:	0.334	train:	0.094	valid:	0.086
iter	360	loss:	0.331	train:	0.093	valid:	0.085
iter	380	loss:	0.329	train:	0.092	valid:	0.085
iter	400	loss:	0.327	train:	0.092	valid:	0.084
iter	420	loss:	0.325	train:	0.091	valid:	0.084
iter	440	loss:	0.323	train:	0.090	valid:	0.084
iter	460	loss:	0.321	train:	0.090	valid:	0.084
iter	480	loss:	0.320	train:	0.089	valid:	0.083
iter	500	loss:	0.318	train:	0.089	valid:	0.083
iter	520	loss:	0.317	train:	0.088	valid:	0.082
iter	540	loss:	0.315	train:	0.088	valid:	0.082
iter	560	loss:	0.314	train:	0.088	valid:	0.082
iter	580	loss:	0.313	train:	0.087	valid:	0.082
iter	600	loss:	0.312	train:	0.087	valid:	0.081
iter	620	loss:	0.311	train:	0.087	valid:	0.081
iter	640	loss:	0.309	train:	0.086	valid:	0.081
iter	660	loss:	0.308	train:	0.086	valid:	0.080
iter	680	loss:	0.307	train:	0.086	valid:	0.080
iter	700	loss:	0.307	train:	0.086	valid:	0.080
iter	720	loss:	0.306	train:	0.085	valid:	0.080
iter	740	loss:	0.305	train:	0.085	valid:	0.080
iter	760	loss:	0.304	train:	0.085	valid:	0.080
iter	780	loss:	0.303	train:	0.085	valid:	0.079
iter	800	loss:	0.302	train:	0.084	valid:	0.079
iter	820	loss:	0.302	train:	0.084	valid:	0.079
iter	840	loss:	0.301	train:	0.084	valid:	0.078
iter	860	loss:	0.300	train:	0.084	valid:	0.078
iter	880	loss:	0.300	train:	0.083	valid:	0.079
iter	900	loss:	0.299	train:	0.083	valid:	0.078
iter	920	loss:	0.298	train:	0.083	valid:	0.078
iter	940	loss:	0.298	train:	0.083	valid:	0.078
iter	960	loss:	0.297	train:	0.083	valid:	0.078
iter	980	loss:	0.296	train:	0.083	valid:	0.078
iter	1000	loss:	0.296	train:	0.082	valid:	0.078





```
In [23]: plt.plot(range(20,1+iters,20),train_error_list)
plt.plot(range(20,1+iters,20),valid_error_list)
plt.legend(['train','test'])
plt.xlabel('iters')
plt.ylabel('error rate')
plt.show()

print('Final result: training error rate %.3f testing error rate %.3f'%(
train_error_list[-1],valid_error_list[-1]))
```



Final result: training error rate 0.082 testing error rate 0.078

```

In [24]: # 4.2 feed forward network with 2 hidden layers
# 4.1 Train a feed forward network with 1 hidden layer
# Here we use Cross Entropy Loss with Softmax activation function
W1 = np.random.normal(0,0.1,(784,32))
W2 = np.random.normal(0,0.1,(32,10))
num_point = len(y_train)
max_iterations = 1000
learning_rate = 1e-5
threshold = 0.0001
loss = num_point
iters = 0
loss_list = []
train_error_list = []
valid_error_list = []
while loss > threshold and iters < max_iterations:
    layer_output_1 = x_train_vector.dot(W1)
    layer_output_2 = layer_output_1.dot(W2)
    y_predict = softmax(layer_output_2)
    loss = - np.sum(y_train * np.log(np.clip(y_predict,1e-15,np.inf)))/num_point
    grad_2 = y_predict - y_train
    W_grad_2 = layer_output_1.T.dot(grad_2)
    grad_1 = (y_predict - y_train).dot(W2.T)
    W_grad_1 = x_train_vector.T.dot(grad_1)


    W2 -= learning_rate*W_grad_2#/num_point
    W1 -= learning_rate*W_grad_1#/num_point

    iters += 1
    loss_list.extend([loss])

    if iters%20 == 0:
        y_max = np.argmax(y_predict,axis=1)
        train_error = 1-np.sum(y_train[range(len(y_train)),y_max])/len(y_train)
        train_error_list.extend([train_error])
        layer_output_1 = x_valid_vector.dot(W1)
        layer_output_2 = layer_output_1.dot(W2)
        y_predict = softmax(layer_output_2)
        y_max = np.argmax(y_predict,axis=1)
        valid_error = 1-np.sum(y_valid[range(len(y_valid)),y_max])/len(y_valid)
        valid_error_list.extend([valid_error])
        print('iter %4d loss: %.3f train: %.3f valid: %.3f'%(iters,loss,train_error,valid_error))
plt.plot(range(1,1+iters),loss_list)
plt.xlabel('iters')
plt.ylabel('loss')
plt.show()


```

```
iter  20 loss: 0.598 train: 0.160 valid: 0.139
iter  40 loss: 0.447 train: 0.123 valid: 0.113
iter  60 loss: 0.398 train: 0.111 valid: 0.103
iter  80 loss: 0.371 train: 0.105 valid: 0.098
iter 100 loss: 0.355 train: 0.101 valid: 0.094
iter 120 loss: 0.343 train: 0.098 valid: 0.090
iter 140 loss: 0.335 train: 0.095 valid: 0.087
iter 160 loss: 0.328 train: 0.093 valid: 0.086
iter 180 loss: 0.323 train: 0.092 valid: 0.085
iter 200 loss: 0.318 train: 0.090 valid: 0.084
iter 220 loss: 0.314 train: 0.089 valid: 0.083
iter 240 loss: 0.311 train: 0.088 valid: 0.082
iter 260 loss: 0.308 train: 0.087 valid: 0.082
iter 280 loss: 0.305 train: 0.087 valid: 0.082
iter 300 loss: 0.303 train: 0.086 valid: 0.081
iter 320 loss: 0.301 train: 0.085 valid: 0.080
iter 340 loss: 0.299 train: 0.084 valid: 0.080
iter 360 loss: 0.297 train: 0.083 valid: 0.080
iter 380 loss: 0.295 train: 0.083 valid: 0.079
iter 400 loss: 0.293 train: 0.082 valid: 0.079
iter 420 loss: 0.292 train: 0.082 valid: 0.079
iter 440 loss: 0.291 train: 0.082 valid: 0.079
iter 460 loss: 0.289 train: 0.081 valid: 0.078
iter 480 loss: 0.288 train: 0.080 valid: 0.077
iter 500 loss: 0.287 train: 0.080 valid: 0.077
iter 520 loss: 0.286 train: 0.080 valid: 0.077
iter 540 loss: 0.284 train: 0.079 valid: 0.076
iter 560 loss: 0.283 train: 0.078 valid: 0.076
iter 580 loss: 0.282 train: 0.078 valid: 0.076
iter 600 loss: 0.281 train: 0.078 valid: 0.076
iter 620 loss: 0.281 train: 0.077 valid: 0.075
iter 640 loss: 0.280 train: 0.077 valid: 0.075
iter 660 loss: 0.279 train: 0.077 valid: 0.075
iter 680 loss: 0.278 train: 0.077 valid: 0.074
iter 700 loss: 0.279 train: 0.078 valid: 0.076
iter 720 loss: 0.304 train: 0.087 valid: 0.085
iter 740 loss: 0.276 train: 0.076 valid: 0.074
iter 760 loss: 0.276 train: 0.077 valid: 0.074
iter 780 loss: 0.295 train: 0.084 valid: 0.082
iter 800 loss: 0.275 train: 0.076 valid: 0.074
iter 820 loss: 0.273 train: 0.075 valid: 0.074
iter 840 loss: 0.273 train: 0.075 valid: 0.074
iter 860 loss: 0.272 train: 0.075 valid: 0.074
iter 880 loss: 0.271 train: 0.075 valid: 0.073
iter 900 loss: 0.271 train: 0.074 valid: 0.073
iter 920 loss: 0.270 train: 0.074 valid: 0.073
iter 940 loss: 0.270 train: 0.074 valid: 0.073
iter 960 loss: 0.270 train: 0.074 valid: 0.073
iter 980 loss: 0.274 train: 0.076 valid: 0.073
iter 1000 loss: 0.274 train: 0.076 valid: 0.073
```



```
In [32]: plt.plot(range(20,1+iters,20),train_error_list)
plt.plot(range(20,1+iters,20),valid_error_list)
plt.legend(['train','test'])
plt.xlabel('iters')
plt.ylabel('error rate')
plt.show()

print('Final result: training error rate %.3f testing error rate %.3f'%(
train_error_list[-1],valid_error_list[-1]))
```




Final result: training error rate 0.076 testing error rate 0.073

```

In [33]: # 4.3 Retrain the network in part 1 with regularization and momentum and
         report the results
W = np.random.normal(0,0.1,(784,10))
num_point = len(y_train)
max_iterations = 1000
learning_rate = 1e-5
threshold = 0.0001
sigma = 1e-5
beta = 0.9
loss = num_point
iters = 0
W_Momentum = np.zeros(W.shape)
loss_list = []
train_error_list = []
valid_error_list = []
while loss > threshold and iters < max_iterations:
    layer_output = x_train_vector.dot(W)
    y_predict = softmax(layer_output)
    loss = - np.sum(y_train * np.log(y_predict))/num_point + sigma*np.su
m(W**2)
    grad = y_predict - y_train
    W_grad = x_train_vector.T.dot(grad) + sigma*2*W
    W_Momentum = beta*W_Momentum + (1-beta)*W_grad
    W -= learning_rate*W_Momentum
    iters += 1
    loss_list.extend([loss])
    if iters%20 == 0:
        y_max = np.argmax(y_predict,axis=1)
        train_error = 1-np.sum(y_train[range(len(y_train)),y_max])/len(y
_train)
        train_error_list.extend([train_error])
        layer_output = x_valid_vector.dot(W)
        y_predict = softmax(layer_output)
        y_max = np.argmax(y_predict,axis=1)
        valid_error = 1-np.sum(y_valid[range(len(y_valid)),y_max])/len(y
_valid)
        valid_error_list.extend([valid_error])
        print('iter %4d loss: %.3f train: %.3f valid: %.3f'%(iters,loss,
train_error,valid_error))
plt.plot(range(1,1+iters),loss_list)
plt.xlabel('iters')
plt.ylabel('loss')
plt.show()


```

iter	20	loss:	0.690	train:	0.196	valid:	0.164
iter	40	loss:	0.496	train:	0.144	valid:	0.125
iter	60	loss:	0.444	train:	0.126	valid:	0.112
iter	80	loss:	0.418	train:	0.117	valid:	0.106
iter	100	loss:	0.401	train:	0.113	valid:	0.102
iter	120	loss:	0.389	train:	0.110	valid:	0.098
iter	140	loss:	0.379	train:	0.107	valid:	0.096
iter	160	loss:	0.371	train:	0.105	valid:	0.094
iter	180	loss:	0.365	train:	0.103	valid:	0.093
iter	200	loss:	0.359	train:	0.102	valid:	0.092
iter	220	loss:	0.354	train:	0.100	valid:	0.091
iter	240	loss:	0.350	train:	0.098	valid:	0.090
iter	260	loss:	0.346	train:	0.097	valid:	0.089
iter	280	loss:	0.342	train:	0.096	valid:	0.088
iter	300	loss:	0.339	train:	0.095	valid:	0.087
iter	320	loss:	0.336	train:	0.094	valid:	0.087
iter	340	loss:	0.334	train:	0.093	valid:	0.086
iter	360	loss:	0.331	train:	0.092	valid:	0.086
iter	380	loss:	0.329	train:	0.092	valid:	0.085
iter	400	loss:	0.327	train:	0.091	valid:	0.085
iter	420	loss:	0.325	train:	0.091	valid:	0.085
iter	440	loss:	0.323	train:	0.090	valid:	0.084
iter	460	loss:	0.322	train:	0.089	valid:	0.083
iter	480	loss:	0.320	train:	0.089	valid:	0.083
iter	500	loss:	0.319	train:	0.088	valid:	0.083
iter	520	loss:	0.317	train:	0.088	valid:	0.082
iter	540	loss:	0.316	train:	0.087	valid:	0.082
iter	560	loss:	0.315	train:	0.087	valid:	0.082
iter	580	loss:	0.314	train:	0.087	valid:	0.081
iter	600	loss:	0.312	train:	0.086	valid:	0.081
iter	620	loss:	0.311	train:	0.086	valid:	0.081
iter	640	loss:	0.310	train:	0.086	valid:	0.080
iter	660	loss:	0.309	train:	0.086	valid:	0.081
iter	680	loss:	0.308	train:	0.085	valid:	0.081
iter	700	loss:	0.307	train:	0.085	valid:	0.080
iter	720	loss:	0.307	train:	0.085	valid:	0.080
iter	740	loss:	0.306	train:	0.085	valid:	0.079
iter	760	loss:	0.305	train:	0.085	valid:	0.079
iter	780	loss:	0.304	train:	0.084	valid:	0.079
iter	800	loss:	0.303	train:	0.084	valid:	0.079
iter	820	loss:	0.303	train:	0.084	valid:	0.078
iter	840	loss:	0.302	train:	0.084	valid:	0.078
iter	860	loss:	0.301	train:	0.084	valid:	0.078
iter	880	loss:	0.301	train:	0.083	valid:	0.078
iter	900	loss:	0.300	train:	0.083	valid:	0.078
iter	920	loss:	0.299	train:	0.083	valid:	0.077
iter	940	loss:	0.299	train:	0.082	valid:	0.077
iter	960	loss:	0.298	train:	0.082	valid:	0.077
iter	980	loss:	0.298	train:	0.082	valid:	0.077
iter	1000	loss:	0.297	train:	0.082	valid:	0.077



```
In [34]: plt.plot(range(20,1+iters,20),train_error_list)
plt.plot(range(20,1+iters,20),valid_error_list)
plt.legend(['train','test'])
plt.xlabel('iters')
plt.ylabel('error rate')
plt.show()

print('Final result: training error rate %.3f testing error rate %.3f'%(
train_error_list[-1],valid_error_list[-1]))
```



Final result: training error rate 0.082 testing error rate 0.077

## 5 Convolutional Neural Network

```
In [23]: from keras.datasets import cifar10
import keras
import numpy as np
import cv2
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train_VGG = np.array(x_train, dtype=np.float32)
x_test_VGG = np.array(x_test, dtype=np.float32)

'''
x_train_VGG = np.zeros((50000, 224, 224, 3), dtype = np.float32)
for index in range(50000):
    x_train_VGG[index, :, :, :] = cv2.resize(x_train[index, :, :, :], (224, 224))
x_test_VGG = np.zeros((10000, 224, 224, 3), dtype = np.float32)
for index in range(10000):
    x_test_VGG[index, :, :, :] = cv2.resize(x_test[index, :, :, :], (224, 224))
'''
y_train_VGG = keras.utils.to_categorical(y_train, 10)
y_test_VGG = keras.utils.to_categorical(y_test, 10)
```

1. Stochastic Gradient Descent [4]
2. Batch Normalization
3. Replace the fully connected layer by average pooling layer
4. Adaptive Gradient [bonus + 2] [3]
5. Nesterov's Accelerated Gradient [bonus + 2] [4]
6. RMSprop [bonus +2] [3]



```

In [151]: import keras
from keras.applications.vgg16 import VGG16
from keras.preprocessing import image
from keras.applications.vgg16 import preprocess_input
from keras.models import Model
from keras.layers import Input, Dense, GlobalAveragePooling2D, Conv2D, MaxPooling2D, Flatten
import numpy as np

img_input = Input((32,32,3))
# Block 1
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv1')(img_input)
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv2')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block1_pool')(x)

# Block 2
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv1')(x)
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv2')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block2_pool')(x)

# Block 3
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv1')(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv2')(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block3_pool')(x)
'''

# Block 4
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block4_pool')(x)

# Block 5
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)
'''

# Final Block
x = Flatten()(x)
x = Dense(128, activation='relu')(x)
x = Dense(32, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

```

```
model = Model(inputs=img_input, outputs=predictions)
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
input_51 (InputLayer)	(None, 32, 32, 3)	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
flatten_11 (Flatten)	(None, 4096)	0
dense_130 (Dense)	(None, 128)	524416
dense_131 (Dense)	(None, 32)	4128
dense_132 (Dense)	(None, 10)	330
=====		
Total params: 2,264,362		
Trainable params: 2,264,362		
Non-trainable params: 0		

```
In [91]: BATCH_SIZE = 128
MAX_EPOCH = 50
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.SGD(lr=1e-3),
              metrics=['accuracy'])
history_SGD_1 = model.fit(x_train_VGG, y_train_VGG,
                          batch_size=BATCH_SIZE,
                          epochs=MAX_EPOCH,
                          verbose=1,
                          validation_data=(x_test_VGG, y_test_VGG))
score = model.evaluate(x_test_VGG, y_test_VGG, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 50000 samples, validate on 10000 samples

Epoch 1/50  
50000/50000 [=====] - 12s 248us/step - loss: 1.9752 - acc: 0.2917 - val\_loss: 1.7017 - val\_acc: 0.3904

Epoch 2/50  
50000/50000 [=====] - 11s 214us/step - loss: 1.6618 - acc: 0.4039 - val\_loss: 1.6621 - val\_acc: 0.4201

Epoch 3/50  
50000/50000 [=====] - 10s 209us/step - loss: 1.5181 - acc: 0.4576 - val\_loss: 1.4824 - val\_acc: 0.4653

Epoch 4/50  
50000/50000 [=====] - 10s 209us/step - loss: 1.4184 - acc: 0.4944 - val\_loss: 1.4592 - val\_acc: 0.4728

Epoch 5/50  
50000/50000 [=====] - 10s 204us/step - loss: 1.3378 - acc: 0.5233 - val\_loss: 1.2845 - val\_acc: 0.5382

Epoch 6/50  
50000/50000 [=====] - 10s 207us/step - loss: 1.2664 - acc: 0.5471 - val\_loss: 1.2978 - val\_acc: 0.5329

Epoch 7/50  
50000/50000 [=====] - 11s 214us/step - loss: 1.2033 - acc: 0.5744 - val\_loss: 1.2471 - val\_acc: 0.5566

Epoch 8/50  
50000/50000 [=====] - 11s 210us/step - loss: 1.1504 - acc: 0.5936 - val\_loss: 1.1983 - val\_acc: 0.5732

Epoch 9/50  
50000/50000 [=====] - 10s 205us/step - loss: 1.1069 - acc: 0.6085 - val\_loss: 1.2177 - val\_acc: 0.5662

Epoch 10/50  
50000/50000 [=====] - 10s 208us/step - loss: 1.0587 - acc: 0.6275 - val\_loss: 1.2358 - val\_acc: 0.5656

Epoch 11/50  
50000/50000 [=====] - 11s 213us/step - loss: 1.0218 - acc: 0.6406 - val\_loss: 1.1479 - val\_acc: 0.6039

Epoch 12/50  
50000/50000 [=====] - 10s 209us/step - loss: 0.9738 - acc: 0.6582 - val\_loss: 1.0958 - val\_acc: 0.6182

Epoch 13/50  
50000/50000 [=====] - 10s 208us/step - loss: 0.9376 - acc: 0.6718 - val\_loss: 1.0436 - val\_acc: 0.6321

Epoch 14/50  
50000/50000 [=====] - 11s 213us/step - loss: 0.9013 - acc: 0.6846 - val\_loss: 1.0632 - val\_acc: 0.6302

Epoch 15/50  
50000/50000 [=====] - 11s 211us/step - loss: 0.8587 - acc: 0.6992 - val\_loss: 1.0733 - val\_acc: 0.6302

Epoch 16/50  
50000/50000 [=====] - 10s 207us/step - loss: 0.8306 - acc: 0.7105 - val\_loss: 1.0912 - val\_acc: 0.6220

Epoch 17/50  
50000/50000 [=====] - 10s 207us/step - loss: 0.7961 - acc: 0.7233 - val\_loss: 0.9963 - val\_acc: 0.6534

Epoch 18/50  
50000/50000 [=====] - 10s 202us/step - loss: 0.7642 - acc: 0.7347 - val\_loss: 1.0398 - val\_acc: 0.6399



Epoch 19/50  
50000/50000 [=====] - 10s 210us/step - loss:

0.7348 - acc: 0.7444 - val\_loss: 1.1518 - val\_acc: 0.6165  
Epoch 20/50  
50000/50000 [=====] - 11s 214us/step - loss:  
0.7021 - acc: 0.7574 - val\_loss: 0.9785 - val\_acc: 0.6661  
Epoch 21/50  
50000/50000 [=====] - 10s 208us/step - loss:  
0.6703 - acc: 0.7667 - val\_loss: 0.9700 - val\_acc: 0.6702  
Epoch 22/50  
50000/50000 [=====] - 10s 208us/step - loss:  
0.6400 - acc: 0.7787 - val\_loss: 1.0687 - val\_acc: 0.6390  
Epoch 23/50  
50000/50000 [=====] - 10s 205us/step - loss:  
0.6118 - acc: 0.7863 - val\_loss: 1.0030 - val\_acc: 0.6695  
Epoch 24/50  
50000/50000 [=====] - 10s 210us/step - loss:  
0.5813 - acc: 0.7986 - val\_loss: 1.0456 - val\_acc: 0.6564  
Epoch 25/50  
50000/50000 [=====] - 11s 214us/step - loss:  
0.5497 - acc: 0.8090 - val\_loss: 1.0500 - val\_acc: 0.6632  
Epoch 26/50  
50000/50000 [=====] - 10s 208us/step - loss:  
0.5239 - acc: 0.8178 - val\_loss: 1.0825 - val\_acc: 0.6642  
Epoch 27/50  
50000/50000 [=====] - 10s 210us/step - loss:  
0.4946 - acc: 0.8297 - val\_loss: 1.1117 - val\_acc: 0.6559  
Epoch 28/50  
50000/50000 [=====] - 10s 203us/step - loss:  
0.4739 - acc: 0.8355 - val\_loss: 1.0388 - val\_acc: 0.6728  
Epoch 29/50  
50000/50000 [=====] - 10s 204us/step - loss:  
0.4409 - acc: 0.8470 - val\_loss: 1.2532 - val\_acc: 0.6253  
Epoch 30/50  
50000/50000 [=====] - 10s 207us/step - loss:  
0.4064 - acc: 0.8620 - val\_loss: 1.1386 - val\_acc: 0.6535  
Epoch 31/50  
50000/50000 [=====] - 11s 216us/step - loss:  
0.3856 - acc: 0.8686 - val\_loss: 1.1013 - val\_acc: 0.6678  
Epoch 32/50  
50000/50000 [=====] - 11s 210us/step - loss:  
0.3565 - acc: 0.8799 - val\_loss: 1.1497 - val\_acc: 0.6546  
Epoch 33/50  
50000/50000 [=====] - 11s 212us/step - loss:  
0.3298 - acc: 0.8874 - val\_loss: 1.1369 - val\_acc: 0.6718  
Epoch 34/50  
50000/50000 [=====] - 11s 212us/step - loss:  
0.3007 - acc: 0.9009 - val\_loss: 1.2001 - val\_acc: 0.6687  
Epoch 35/50  
50000/50000 [=====] - 11s 212us/step - loss:  
0.2783 - acc: 0.9100 - val\_loss: 1.2343 - val\_acc: 0.6626  
Epoch 36/50  
50000/50000 [=====] - 10s 205us/step - loss:  
0.2471 - acc: 0.9196 - val\_loss: 1.3716 - val\_acc: 0.6495  
Epoch 37/50  
50000/50000 [=====] - 10s 204us/step - loss:  
0.2304 - acc: 0.9258 - val\_loss: 1.4500 - val\_acc: 0.6355  
Epoch 38/50  
50000/50000 [=====] - 10s 204us/step - loss:

```
0.2255 - acc: 0.9318 - val_loss: 1.2501 - val_acc: 0.6774
Epoch 39/50
50000/50000 [=====] - 10s 204us/step - loss:
0.1734 - acc: 0.9472 - val_loss: 1.2839 - val_acc: 0.6765
Epoch 40/50
50000/50000 [=====] - 11s 211us/step - loss:
0.1560 - acc: 0.9543 - val_loss: 1.3359 - val_acc: 0.6747
Epoch 41/50
50000/50000 [=====] - 10s 205us/step - loss:
0.1806 - acc: 0.9535 - val_loss: 1.3756 - val_acc: 0.6714
Epoch 42/50
50000/50000 [=====] - 10s 204us/step - loss:
0.1670 - acc: 0.9565 - val_loss: 1.3529 - val_acc: 0.6800
Epoch 43/50
50000/50000 [=====] - 10s 205us/step - loss:
0.0889 - acc: 0.9795 - val_loss: 1.4415 - val_acc: 0.6755
Epoch 44/50
50000/50000 [=====] - 11s 210us/step - loss:
0.0670 - acc: 0.9865 - val_loss: 1.5803 - val_acc: 0.6643
Epoch 45/50
50000/50000 [=====] - 10s 202us/step - loss:
0.0526 - acc: 0.9910 - val_loss: 1.5749 - val_acc: 0.6721
Epoch 46/50
50000/50000 [=====] - 11s 216us/step - loss:
0.1306 - acc: 0.9731 - val_loss: 1.4522 - val_acc: 0.6661
Epoch 47/50
50000/50000 [=====] - 10s 205us/step - loss:
0.0436 - acc: 0.9944 - val_loss: 1.5846 - val_acc: 0.6812
Epoch 48/50
50000/50000 [=====] - 10s 203us/step - loss:
0.0259 - acc: 0.9981 - val_loss: 1.6626 - val_acc: 0.6796
Epoch 49/50
50000/50000 [=====] - 11s 215us/step - loss:
0.0198 - acc: 0.9990 - val_loss: 1.6924 - val_acc: 0.6787
Epoch 50/50
50000/50000 [=====] - 11s 210us/step - loss:
0.0160 - acc: 0.9995 - val_loss: 1.7450 - val_acc: 0.6794
Test loss: 1.7449563205718994
Test accuracy: 0.6794
```

```
In [168]: # summarize history for loss
plt.plot(history_SGD_1.history['loss'])
plt.plot(history_SGD_1.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.ylim(0,2.2)
plt.legend(['train', 'test'], loc='upper left')
plt.show()

# summarize history for accuracy
plt.plot(history_SGD_1.history['acc'])
plt.plot(history_SGD_1.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.ylim(0.1,1)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```

In [149]: import keras
from keras.applications.vgg16 import VGG16
from keras.preprocessing import image
from keras.applications.vgg16 import preprocess_input
from keras.models import Model
from keras.layers import Input, Dense, GlobalAveragePooling2D, Conv2D, MaxPooling2D, Flatten
import numpy as np

img_input = Input((32,32,3))
# Block 1
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv1')(img_input)
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv2')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block1_pool')(x)

# Block 2
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv1')(x)
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv2')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block2_pool')(x)

# Block 3
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv1')(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv2')(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block3_pool')(x)
'''

# Block 4
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block4_pool')(x)

# Block 5
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)
'''

# Final Block
x = GlobalAveragePooling2D()(x)
x = Dense(128, activation='relu')(x)
x = Dense(32, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

```



```
model = Model(inputs=img_input, outputs=predictions)
model.summary()
```

Layer (type)	Output Shape	Param #
input_50 (InputLayer)	(None, 32, 32, 3)	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
global_average_pooling2d_36	(None, 256)	0
dense_127 (Dense)	(None, 128)	32896
dense_128 (Dense)	(None, 32)	4128
dense_129 (Dense)	(None, 10)	330
Total params: 1,772,842		
Trainable params: 1,772,842		
Non-trainable params: 0		

```
In [93]: BATCH_SIZE = 128
MAX_EPOCH = 50
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.SGD(lr=1e-3),
              metrics=['accuracy'])
history_SGD_2 = model.fit(x_train_VGG, y_train_VGG,
                          batch_size=BATCH_SIZE,
                          epochs=MAX_EPOCH,
                          verbose=1,
                          validation_data=(x_test_VGG, y_test_VGG))
score = model.evaluate(x_test_VGG, y_test_VGG, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 50000 samples, validate on 10000 samples

Epoch 1/50  
50000/50000 [=====] - 12s 236us/step - loss: 2.1698 - acc: 0.1813 - val\_loss: 2.0338 - val\_acc: 0.2514

Epoch 2/50  
50000/50000 [=====] - 10s 191us/step - loss: 1.9871 - acc: 0.2826 - val\_loss: 2.0732 - val\_acc: 0.2518

Epoch 3/50  
50000/50000 [=====] - 10s 194us/step - loss: 1.8747 - acc: 0.3281 - val\_loss: 1.8847 - val\_acc: 0.3191

Epoch 4/50  
50000/50000 [=====] - 9s 189us/step - loss: 1.7769 - acc: 0.3587 - val\_loss: 1.7226 - val\_acc: 0.3710

Epoch 5/50  
50000/50000 [=====] - 10s 190us/step - loss: 1.7038 - acc: 0.3837 - val\_loss: 1.6612 - val\_acc: 0.3947

Epoch 6/50  
50000/50000 [=====] - 10s 198us/step - loss: 1.6365 - acc: 0.4097 - val\_loss: 1.5955 - val\_acc: 0.4206

Epoch 7/50  
50000/50000 [=====] - 10s 193us/step - loss: 1.5869 - acc: 0.4264 - val\_loss: 1.5297 - val\_acc: 0.4540

Epoch 8/50  
50000/50000 [=====] - 10s 192us/step - loss: 1.5411 - acc: 0.4420 - val\_loss: 1.4800 - val\_acc: 0.4619

Epoch 9/50  
50000/50000 [=====] - 10s 193us/step - loss: 1.4897 - acc: 0.4628 - val\_loss: 1.6050 - val\_acc: 0.4322

Epoch 10/50  
50000/50000 [=====] - 10s 197us/step - loss: 1.4610 - acc: 0.4739 - val\_loss: 1.4712 - val\_acc: 0.4629

Epoch 11/50  
50000/50000 [=====] - 10s 198us/step - loss: 1.4227 - acc: 0.4861 - val\_loss: 1.4915 - val\_acc: 0.4540

Epoch 12/50  
50000/50000 [=====] - 10s 193us/step - loss: 1.3902 - acc: 0.5000 - val\_loss: 1.3962 - val\_acc: 0.4965

Epoch 13/50  
50000/50000 [=====] - 9s 185us/step - loss: 1.3535 - acc: 0.5115 - val\_loss: 1.4236 - val\_acc: 0.4827

Epoch 14/50  
50000/50000 [=====] - 10s 196us/step - loss: 1.3242 - acc: 0.5251 - val\_loss: 1.2781 - val\_acc: 0.5440

Epoch 15/50  
50000/50000 [=====] - 10s 193us/step - loss: 1.2976 - acc: 0.5351 - val\_loss: 1.3272 - val\_acc: 0.5198

Epoch 16/50  
50000/50000 [=====] - 10s 194us/step - loss: 1.2688 - acc: 0.5468 - val\_loss: 1.4544 - val\_acc: 0.4623

Epoch 17/50  
50000/50000 [=====] - 10s 191us/step - loss: 1.2383 - acc: 0.5593 - val\_loss: 1.2560 - val\_acc: 0.5571

Epoch 18/50  
50000/50000 [=====] - 10s 195us/step - loss: 1.2151 - acc: 0.5681 - val\_loss: 1.2006 - val\_acc: 0.5765



Epoch 19/50  
50000/50000 [=====] - 10s 193us/step - loss:

```
1.1941 - acc: 0.5755 - val_loss: 1.1932 - val_acc: 0.5758
Epoch 20/50
50000/50000 [=====] - 10s 192us/step - loss:
1.1632 - acc: 0.5904 - val_loss: 1.1541 - val_acc: 0.5875
Epoch 21/50
50000/50000 [=====] - 10s 192us/step - loss:
1.1356 - acc: 0.5994 - val_loss: 1.1136 - val_acc: 0.6071
Epoch 22/50
50000/50000 [=====] - 9s 187us/step - loss: 1.
1146 - acc: 0.6075 - val_loss: 1.1093 - val_acc: 0.6089
Epoch 23/50
50000/50000 [=====] - 9s 189us/step - loss: 1.
0826 - acc: 0.6171 - val_loss: 1.1850 - val_acc: 0.5690
Epoch 24/50
50000/50000 [=====] - 10s 197us/step - loss:
1.0665 - acc: 0.6263 - val_loss: 1.2511 - val_acc: 0.5557
Epoch 25/50
50000/50000 [=====] - 10s 194us/step - loss:
1.0422 - acc: 0.6370 - val_loss: 1.1241 - val_acc: 0.6054
Epoch 26/50
50000/50000 [=====] - 10s 199us/step - loss:
1.0199 - acc: 0.6433 - val_loss: 1.0213 - val_acc: 0.6413
Epoch 27/50
50000/50000 [=====] - 10s 194us/step - loss:
0.9993 - acc: 0.6505 - val_loss: 1.0818 - val_acc: 0.6072
Epoch 28/50
50000/50000 [=====] - 10s 196us/step - loss:
0.9805 - acc: 0.6545 - val_loss: 1.0325 - val_acc: 0.6404
Epoch 29/50
50000/50000 [=====] - 10s 196us/step - loss:
0.9595 - acc: 0.6648 - val_loss: 1.0289 - val_acc: 0.6426
Epoch 30/50
50000/50000 [=====] - 10s 200us/step - loss:
0.9464 - acc: 0.6694 - val_loss: 0.9805 - val_acc: 0.6554
Epoch 31/50
50000/50000 [=====] - 10s 196us/step - loss:
0.9257 - acc: 0.6770 - val_loss: 0.9781 - val_acc: 0.6498
Epoch 32/50
50000/50000 [=====] - 10s 199us/step - loss:
0.9087 - acc: 0.6834 - val_loss: 1.0768 - val_acc: 0.6180
Epoch 33/50
50000/50000 [=====] - 10s 195us/step - loss:
0.8891 - acc: 0.6904 - val_loss: 0.9895 - val_acc: 0.6535
Epoch 34/50
50000/50000 [=====] - 10s 201us/step - loss:
0.8799 - acc: 0.6935 - val_loss: 0.9213 - val_acc: 0.6815
Epoch 35/50
50000/50000 [=====] - 10s 192us/step - loss:
0.8611 - acc: 0.6986 - val_loss: 0.9255 - val_acc: 0.6796
Epoch 36/50
50000/50000 [=====] - 10s 193us/step - loss:
0.8421 - acc: 0.7073 - val_loss: 0.9887 - val_acc: 0.6469
Epoch 37/50
50000/50000 [=====] - 10s 195us/step - loss:
0.8322 - acc: 0.7087 - val_loss: 1.0056 - val_acc: 0.6412
Epoch 38/50
50000/50000 [=====] - 10s 195us/step - loss:
```

```
0.8114 - acc: 0.7178 - val_loss: 1.0355 - val_acc: 0.6296
Epoch 39/50
50000/50000 [=====] - 10s 197us/step - loss:
0.7989 - acc: 0.7230 - val_loss: 0.9568 - val_acc: 0.6606
Epoch 40/50
50000/50000 [=====] - 10s 196us/step - loss:
0.7905 - acc: 0.7238 - val_loss: 0.9639 - val_acc: 0.6604
Epoch 41/50
50000/50000 [=====] - 10s 194us/step - loss:
0.7710 - acc: 0.7318 - val_loss: 1.0069 - val_acc: 0.6506
Epoch 42/50
50000/50000 [=====] - 10s 196us/step - loss:
0.7575 - acc: 0.7375 - val_loss: 0.9849 - val_acc: 0.6569
Epoch 43/50
50000/50000 [=====] - 10s 206us/step - loss:
0.7457 - acc: 0.7419 - val_loss: 0.8761 - val_acc: 0.6954
Epoch 44/50
50000/50000 [=====] - 10s 199us/step - loss:
0.7340 - acc: 0.7440 - val_loss: 0.9148 - val_acc: 0.6801
Epoch 45/50
50000/50000 [=====] - 10s 197us/step - loss:
0.7163 - acc: 0.7511 - val_loss: 0.9015 - val_acc: 0.6886
Epoch 46/50
50000/50000 [=====] - 10s 195us/step - loss:
0.7084 - acc: 0.7550 - val_loss: 0.9246 - val_acc: 0.6807
Epoch 47/50
50000/50000 [=====] - 10s 196us/step - loss:
0.6960 - acc: 0.7585 - val_loss: 0.9215 - val_acc: 0.6777
Epoch 48/50
50000/50000 [=====] - 10s 198us/step - loss:
0.6788 - acc: 0.7668 - val_loss: 0.9556 - val_acc: 0.6783
Epoch 49/50
50000/50000 [=====] - 10s 202us/step - loss:
0.6645 - acc: 0.7685 - val_loss: 0.8713 - val_acc: 0.6931
Epoch 50/50
50000/50000 [=====] - 10s 199us/step - loss:
0.6510 - acc: 0.7730 - val_loss: 0.9348 - val_acc: 0.6785
Test loss: 0.9347507353782654
Test accuracy: 0.6785
```

```
In [169]: # summarize history for loss
plt.plot(history_SGD_2.history['loss'])
plt.plot(history_SGD_2.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.ylim(0,2.2)
plt.legend(['train', 'test'], loc='upper left')
plt.show()

# summarize history for accuracy
plt.plot(history_SGD_2.history['acc'])
plt.plot(history_SGD_2.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.ylim(0.1,1)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```

In [120]: import keras
from keras.applications.vgg16 import VGG16
from keras.preprocessing import image
from keras.applications.vgg16 import preprocess_input
from keras.models import Model
from keras.layers import Input, Dense, GlobalAveragePooling2D, Conv2D, MaxPooling2D, BatchNormalization
import numpy as np

img_input = Input((32,32,3))
# Block 1
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv1')(img_input)
x = BatchNormalization()(x)
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv2')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block1_pool')(x)

# Block 2
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv1')(x)
x = BatchNormalization()(x)
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv2')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block2_pool')(x)

# Block 3
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv1')(x)
x = BatchNormalization()(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv2')(x)
x = BatchNormalization()(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv3')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block3_pool')(x)

# Block 4
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block4_pool')(x)

# Block 5
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv3')(x)

```

```
x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)
'''
# Final Block
x = GlobalAveragePooling2D()(x)
x = Dense(128, activation='relu')(x)
x = Dense(32, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

model = Model(inputs=img_input, outputs=predictions)
model.summary()
```



Layer (type)	Output Shape	Param #
input_42 (InputLayer)	(None, 32, 32, 3)	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
batch_normalization_140 (Batch Normalization)	(None, 32, 32, 64)	256
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
batch_normalization_141 (Batch Normalization)	(None, 32, 32, 64)	256
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
batch_normalization_142 (Batch Normalization)	(None, 16, 16, 128)	512
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
batch_normalization_143 (Batch Normalization)	(None, 16, 16, 128)	512
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
batch_normalization_144 (Batch Normalization)	(None, 8, 8, 256)	1024
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
batch_normalization_145 (Batch Normalization)	(None, 8, 8, 256)	1024
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
batch_normalization_146 (Batch Normalization)	(None, 8, 8, 256)	1024
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
global_average_pooling2d_29 (Global Average Pooling)	(None, 256)	0
dense_103 (Dense)	(None, 128)	32896
dense_104 (Dense)	(None, 32)	4128
dense_105 (Dense)	(None, 10)	330
Total params: 1,777,450		
Trainable params: 1,775,146		
Non-trainable params: 2,304		

```
In [95]: BATCH_SIZE = 128
MAX_EPOCH = 50
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.SGD(lr=1e-3),
              metrics=['accuracy'])
history_SGD_3 = model.fit(x_train_VGG, y_train_VGG,
                          batch_size=BATCH_SIZE,
                          epochs=MAX_EPOCH,
                          verbose=1,
                          validation_data=(x_test_VGG, y_test_VGG))
score = model.evaluate(x_test_VGG, y_test_VGG, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 50000 samples, validate on 10000 samples

Epoch 1/50  
50000/50000 [=====] - 13s 267us/step - loss: 2.0653 - acc: 0.2551 - val\_loss: 1.9189 - val\_acc: 0.3186

Epoch 2/50  
50000/50000 [=====] - 11s 212us/step - loss: 1.8270 - acc: 0.3557 - val\_loss: 1.7386 - val\_acc: 0.3887

Epoch 3/50  
50000/50000 [=====] - 11s 219us/step - loss: 1.6504 - acc: 0.4240 - val\_loss: 1.5751 - val\_acc: 0.4490

Epoch 4/50  
50000/50000 [=====] - 11s 217us/step - loss: 1.5010 - acc: 0.4786 - val\_loss: 1.4890 - val\_acc: 0.4786

Epoch 5/50  
50000/50000 [=====] - 11s 216us/step - loss: 1.3882 - acc: 0.5161 - val\_loss: 1.3747 - val\_acc: 0.5193

Epoch 6/50  
50000/50000 [=====] - 11s 215us/step - loss: 1.2998 - acc: 0.5458 - val\_loss: 1.3190 - val\_acc: 0.5288

Epoch 7/50  
50000/50000 [=====] - 11s 214us/step - loss: 1.2288 - acc: 0.5677 - val\_loss: 1.2938 - val\_acc: 0.5375

Epoch 8/50  
50000/50000 [=====] - 11s 220us/step - loss: 1.1690 - acc: 0.5895 - val\_loss: 1.2071 - val\_acc: 0.5718

Epoch 9/50  
50000/50000 [=====] - 11s 217us/step - loss: 1.1165 - acc: 0.6088 - val\_loss: 1.2479 - val\_acc: 0.5607

Epoch 10/50  
50000/50000 [=====] - 11s 220us/step - loss: 1.0692 - acc: 0.6265 - val\_loss: 1.1236 - val\_acc: 0.5972

Epoch 11/50  
50000/50000 [=====] - 11s 216us/step - loss: 1.0271 - acc: 0.6412 - val\_loss: 1.3015 - val\_acc: 0.5448

Epoch 12/50  
50000/50000 [=====] - 11s 213us/step - loss: 0.9874 - acc: 0.6557 - val\_loss: 1.1231 - val\_acc: 0.6048

Epoch 13/50  
50000/50000 [=====] - 11s 213us/step - loss: 0.9529 - acc: 0.6671 - val\_loss: 1.0636 - val\_acc: 0.6225

Epoch 14/50  
50000/50000 [=====] - 11s 223us/step - loss: 0.9190 - acc: 0.6804 - val\_loss: 1.0173 - val\_acc: 0.6410

Epoch 15/50  
50000/50000 [=====] - 11s 218us/step - loss: 0.8892 - acc: 0.6905 - val\_loss: 1.0088 - val\_acc: 0.6483

Epoch 16/50  
50000/50000 [=====] - 11s 216us/step - loss: 0.8590 - acc: 0.7012 - val\_loss: 1.0066 - val\_acc: 0.6490

Epoch 17/50  
50000/50000 [=====] - 11s 216us/step - loss: 0.8309 - acc: 0.7102 - val\_loss: 1.0090 - val\_acc: 0.6415



Epoch 18/50  
50000/50000 [=====] - 11s 216us/step - loss: 0.8043 - acc: 0.7235 - val\_loss: 0.9519 - val\_acc: 0.6694

Epoch 19/50  
50000/50000 [=====] - 11s 218us/step - loss:

```
0.7792 - acc: 0.7324 - val_loss: 1.0536 - val_acc: 0.6307
Epoch 20/50
50000/50000 [=====] - 11s 215us/step - loss:
0.7561 - acc: 0.7388 - val_loss: 0.9626 - val_acc: 0.6612
Epoch 21/50
50000/50000 [=====] - 11s 212us/step - loss:
0.7320 - acc: 0.7467 - val_loss: 1.0048 - val_acc: 0.6516
Epoch 22/50
50000/50000 [=====] - 11s 219us/step - loss:
0.7108 - acc: 0.7548 - val_loss: 0.9563 - val_acc: 0.6602
Epoch 23/50
50000/50000 [=====] - 11s 215us/step - loss:
0.6899 - acc: 0.7635 - val_loss: 0.9713 - val_acc: 0.6639
Epoch 24/50
50000/50000 [=====] - 11s 217us/step - loss:
0.6662 - acc: 0.7728 - val_loss: 0.8973 - val_acc: 0.6836
Epoch 25/50
50000/50000 [=====] - 11s 216us/step - loss:
0.6460 - acc: 0.7812 - val_loss: 1.0141 - val_acc: 0.6505
Epoch 26/50
50000/50000 [=====] - 11s 216us/step - loss:
0.6236 - acc: 0.7894 - val_loss: 0.9356 - val_acc: 0.6736
Epoch 27/50
50000/50000 [=====] - 11s 220us/step - loss:
0.6047 - acc: 0.7960 - val_loss: 0.9178 - val_acc: 0.6801
Epoch 28/50
50000/50000 [=====] - 11s 220us/step - loss:
0.5839 - acc: 0.8058 - val_loss: 0.9160 - val_acc: 0.6842
Epoch 29/50
50000/50000 [=====] - 11s 217us/step - loss:
0.5644 - acc: 0.8126 - val_loss: 0.9283 - val_acc: 0.6798
Epoch 30/50
50000/50000 [=====] - 11s 217us/step - loss:
0.5447 - acc: 0.8198 - val_loss: 0.8938 - val_acc: 0.6874
Epoch 31/50
50000/50000 [=====] - 11s 215us/step - loss:
0.5244 - acc: 0.8265 - val_loss: 0.9177 - val_acc: 0.6859
Epoch 32/50
50000/50000 [=====] - 11s 220us/step - loss:
0.5073 - acc: 0.8341 - val_loss: 0.8572 - val_acc: 0.7034
Epoch 33/50
50000/50000 [=====] - 11s 217us/step - loss:
0.4869 - acc: 0.8424 - val_loss: 0.8759 - val_acc: 0.6992
Epoch 34/50
50000/50000 [=====] - 11s 218us/step - loss:
0.4702 - acc: 0.8486 - val_loss: 0.8441 - val_acc: 0.7114
Epoch 35/50
50000/50000 [=====] - 11s 216us/step - loss:
0.4515 - acc: 0.8554 - val_loss: 0.8544 - val_acc: 0.7033
Epoch 36/50
50000/50000 [=====] - 11s 223us/step - loss:
0.4330 - acc: 0.8630 - val_loss: 0.8765 - val_acc: 0.7006
Epoch 37/50
50000/50000 [=====] - 11s 220us/step - loss:
0.4173 - acc: 0.8695 - val_loss: 0.8517 - val_acc: 0.7099
Epoch 38/50
50000/50000 [=====] - 11s 221us/step - loss:
```

```
0.3999 - acc: 0.8760 - val_loss: 0.8574 - val_acc: 0.7107
Epoch 39/50
50000/50000 [=====] - 11s 222us/step - loss:
0.3808 - acc: 0.8835 - val_loss: 0.8778 - val_acc: 0.7056
Epoch 40/50
50000/50000 [=====] - 11s 222us/step - loss:
0.3657 - acc: 0.8900 - val_loss: 0.8529 - val_acc: 0.7106
Epoch 41/50
50000/50000 [=====] - 11s 217us/step - loss:
0.3491 - acc: 0.8957 - val_loss: 0.8783 - val_acc: 0.7033
Epoch 42/50
50000/50000 [=====] - 11s 225us/step - loss:
0.3342 - acc: 0.9017 - val_loss: 0.8450 - val_acc: 0.7191
Epoch 43/50
50000/50000 [=====] - 11s 214us/step - loss:
0.3185 - acc: 0.9080 - val_loss: 0.8649 - val_acc: 0.7181
Epoch 44/50
50000/50000 [=====] - 11s 219us/step - loss:
0.3010 - acc: 0.9149 - val_loss: 0.8572 - val_acc: 0.7171
Epoch 45/50
50000/50000 [=====] - 11s 218us/step - loss:
0.2869 - acc: 0.9216 - val_loss: 0.8575 - val_acc: 0.7155
Epoch 46/50
50000/50000 [=====] - 11s 218us/step - loss:
0.2725 - acc: 0.9263 - val_loss: 0.8722 - val_acc: 0.7122
Epoch 47/50
50000/50000 [=====] - 11s 215us/step - loss:
0.2588 - acc: 0.9301 - val_loss: 0.8565 - val_acc: 0.7220
Epoch 48/50
50000/50000 [=====] - 11s 219us/step - loss:
0.2442 - acc: 0.9361 - val_loss: 0.8675 - val_acc: 0.7231
Epoch 49/50
50000/50000 [=====] - 11s 214us/step - loss:
0.2325 - acc: 0.9407 - val_loss: 0.9260 - val_acc: 0.7034
Epoch 50/50
50000/50000 [=====] - 11s 215us/step - loss:
0.2180 - acc: 0.9459 - val_loss: 0.9274 - val_acc: 0.7087
Test loss: 0.9273944653511047
Test accuracy: 0.7087
```

```
In [170]: # summarize history for loss
plt.plot(history_SGD_3.history['loss'])
plt.plot(history_SGD_3.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.ylim(0,2.2)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
# summarize history for accuracy
plt.plot(history_SGD_3.history['acc'])
plt.plot(history_SGD_3.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.ylim(0.1,1)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
In [121]: BATCH_SIZE = 128
MAX_EPOCH = 50
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.RMSprop(),
              metrics=['accuracy'])
history_RMSprop_nodrop = model.fit(x_train_VGG, y_train_VGG,
                                  batch_size=BATCH_SIZE,
                                  epochs=MAX_EPOCH,
                                  verbose=1,
                                  validation_data=(x_test_VGG, y_test_VGG))
score = model.evaluate(x_test_VGG, y_test_VGG, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 50000 samples, validate on 10000 samples

Epoch 1/50  
50000/50000 [=====] - 15s 295us/step - loss: 1.3123 - acc: 0.5213 - val\_loss: 1.5995 - val\_acc: 0.4860

Epoch 2/50  
50000/50000 [=====] - 12s 230us/step - loss: 0.7803 - acc: 0.7228 - val\_loss: 0.9352 - val\_acc: 0.6845

Epoch 3/50  
50000/50000 [=====] - 11s 229us/step - loss: 0.5757 - acc: 0.7998 - val\_loss: 1.0337 - val\_acc: 0.6909

Epoch 4/50  
50000/50000 [=====] - 12s 231us/step - loss: 0.4524 - acc: 0.8434 - val\_loss: 1.0933 - val\_acc: 0.6813

Epoch 5/50  
50000/50000 [=====] - 11s 223us/step - loss: 0.3579 - acc: 0.8771 - val\_loss: 0.7822 - val\_acc: 0.7710

Epoch 6/50  
50000/50000 [=====] - 11s 227us/step - loss: 0.2856 - acc: 0.9009 - val\_loss: 0.8344 - val\_acc: 0.7686

Epoch 7/50  
50000/50000 [=====] - 11s 225us/step - loss: 0.2277 - acc: 0.9211 - val\_loss: 0.6931 - val\_acc: 0.8003

Epoch 8/50  
50000/50000 [=====] - 11s 228us/step - loss: 0.1827 - acc: 0.9369 - val\_loss: 0.7239 - val\_acc: 0.8042

Epoch 9/50  
50000/50000 [=====] - 11s 222us/step - loss: 0.1470 - acc: 0.9500 - val\_loss: 0.7129 - val\_acc: 0.8241

Epoch 10/50  
50000/50000 [=====] - 11s 224us/step - loss: 0.1216 - acc: 0.9575 - val\_loss: 0.8054 - val\_acc: 0.8032

Epoch 11/50  
50000/50000 [=====] - 11s 228us/step - loss: 0.1060 - acc: 0.9637 - val\_loss: 0.8918 - val\_acc: 0.7932

Epoch 12/50  
50000/50000 [=====] - 11s 229us/step - loss: 0.0931 - acc: 0.9682 - val\_loss: 0.9365 - val\_acc: 0.7951

Epoch 13/50  
50000/50000 [=====] - 11s 225us/step - loss: 0.0820 - acc: 0.9722 - val\_loss: 0.8201 - val\_acc: 0.8127

Epoch 14/50  
50000/50000 [=====] - 11s 224us/step - loss: 0.0753 - acc: 0.9741 - val\_loss: 0.9258 - val\_acc: 0.8059

Epoch 15/50  
50000/50000 [=====] - 11s 226us/step - loss: 0.0675 - acc: 0.9769 - val\_loss: 0.9355 - val\_acc: 0.8083

Epoch 16/50  
50000/50000 [=====] - 11s 227us/step - loss: 0.0655 - acc: 0.9783 - val\_loss: 0.8589 - val\_acc: 0.8133

Epoch 17/50  
50000/50000 [=====] - 11s 225us/step - loss: 0.0584 - acc: 0.9801 - val\_loss: 0.9247 - val\_acc: 0.8113

Epoch 18/50  
50000/50000 [=====] - 11s 227us/step - loss: 0.0569 - acc: 0.9824 - val\_loss: 1.0206 - val\_acc: 0.8022



Epoch 19/50  
50000/50000 [=====] - 11s 225us/step - loss:



```
0.0508 - acc: 0.9834 - val_loss: 0.9758 - val_acc: 0.8221
Epoch 20/50
50000/50000 [=====] - 11s 227us/step - loss:
0.0490 - acc: 0.9837 - val_loss: 0.8932 - val_acc: 0.8328
Epoch 21/50
50000/50000 [=====] - 12s 230us/step - loss:
0.0462 - acc: 0.9847 - val_loss: 0.9501 - val_acc: 0.8197
Epoch 22/50
50000/50000 [=====] - 11s 226us/step - loss:
0.0427 - acc: 0.9857 - val_loss: 1.0969 - val_acc: 0.8045
Epoch 23/50
50000/50000 [=====] - 11s 221us/step - loss:
0.0431 - acc: 0.9861 - val_loss: 0.8489 - val_acc: 0.8451
Epoch 24/50
50000/50000 [=====] - 11s 223us/step - loss:
0.0394 - acc: 0.9870 - val_loss: 0.8974 - val_acc: 0.8274
Epoch 25/50
50000/50000 [=====] - 11s 224us/step - loss:
0.0395 - acc: 0.9872 - val_loss: 1.0136 - val_acc: 0.8182
Epoch 26/50
50000/50000 [=====] - 11s 221us/step - loss:
0.0371 - acc: 0.9885 - val_loss: 1.0285 - val_acc: 0.8191
Epoch 27/50
50000/50000 [=====] - 11s 227us/step - loss:
0.0338 - acc: 0.9890 - val_loss: 0.9837 - val_acc: 0.8343
Epoch 28/50
50000/50000 [=====] - 11s 227us/step - loss:
0.0350 - acc: 0.9888 - val_loss: 0.9214 - val_acc: 0.8454
Epoch 29/50
50000/50000 [=====] - 11s 225us/step - loss:
0.0357 - acc: 0.9892 - val_loss: 1.0120 - val_acc: 0.8291
Epoch 30/50
50000/50000 [=====] - 11s 228us/step - loss:
0.0338 - acc: 0.9897 - val_loss: 1.3033 - val_acc: 0.7920
Epoch 31/50
50000/50000 [=====] - 11s 223us/step - loss:
0.0328 - acc: 0.9895 - val_loss: 0.9596 - val_acc: 0.8416
Epoch 32/50
50000/50000 [=====] - 11s 225us/step - loss:
0.0303 - acc: 0.9907 - val_loss: 0.9460 - val_acc: 0.8380
Epoch 33/50
50000/50000 [=====] - 11s 230us/step - loss:
0.0321 - acc: 0.9904 - val_loss: 0.8325 - val_acc: 0.8437
Epoch 34/50
50000/50000 [=====] - 11s 230us/step - loss:
0.0284 - acc: 0.9910 - val_loss: 0.9570 - val_acc: 0.8408
Epoch 35/50
50000/50000 [=====] - 11s 229us/step - loss:
0.0279 - acc: 0.9906 - val_loss: 1.1044 - val_acc: 0.8321
Epoch 36/50
50000/50000 [=====] - 12s 233us/step - loss:
0.0279 - acc: 0.9916 - val_loss: 0.9772 - val_acc: 0.8456
Epoch 37/50
50000/50000 [=====] - 11s 227us/step - loss:
0.0271 - acc: 0.9919 - val_loss: 1.1696 - val_acc: 0.8348
Epoch 38/50
50000/50000 [=====] - 11s 226us/step - loss:
```

```
0.0290 - acc: 0.9911 - val_loss: 0.9907 - val_acc: 0.8447
Epoch 39/50
50000/50000 [=====] - 11s 229us/step - loss:
0.0260 - acc: 0.9923 - val_loss: 1.0360 - val_acc: 0.8289
Epoch 40/50
50000/50000 [=====] - 12s 235us/step - loss:
0.0245 - acc: 0.9927 - val_loss: 1.1326 - val_acc: 0.8315
Epoch 41/50
50000/50000 [=====] - 11s 225us/step - loss:
0.0254 - acc: 0.9923 - val_loss: 1.2312 - val_acc: 0.8142
Epoch 42/50
50000/50000 [=====] - 11s 224us/step - loss:
0.0246 - acc: 0.9923 - val_loss: 1.1483 - val_acc: 0.8357
Epoch 43/50
50000/50000 [=====] - 12s 231us/step - loss:
0.0263 - acc: 0.9923 - val_loss: 1.0552 - val_acc: 0.8317
Epoch 44/50
50000/50000 [=====] - 11s 229us/step - loss:
0.0238 - acc: 0.9929 - val_loss: 1.0707 - val_acc: 0.8396
Epoch 45/50
50000/50000 [=====] - 11s 227us/step - loss:
0.0236 - acc: 0.9932 - val_loss: 0.9883 - val_acc: 0.8398
Epoch 46/50
50000/50000 [=====] - 11s 228us/step - loss:
0.0220 - acc: 0.9934 - val_loss: 1.2171 - val_acc: 0.8219
Epoch 47/50
50000/50000 [=====] - 12s 230us/step - loss:
0.0241 - acc: 0.9926 - val_loss: 1.2704 - val_acc: 0.8359
Epoch 48/50
50000/50000 [=====] - 11s 230us/step - loss:
0.0221 - acc: 0.9933 - val_loss: 1.1380 - val_acc: 0.8394
Epoch 49/50
50000/50000 [=====] - 11s 224us/step - loss:
0.0209 - acc: 0.9938 - val_loss: 1.0960 - val_acc: 0.8363
Epoch 50/50
50000/50000 [=====] - 11s 225us/step - loss:
0.0213 - acc: 0.9936 - val_loss: 1.0131 - val_acc: 0.8341
Test loss: 1.0131436536431313
Test accuracy: 0.8341
```

```
In [171]: # summarize history for loss
plt.plot(history_RMSprop_nodrop.history['loss'])
plt.plot(history_RMSprop_nodrop.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.ylim(0,2.2)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
# summarize history for accuracy
plt.plot(history_RMSprop_nodrop.history['acc'])
plt.plot(history_RMSprop_nodrop.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.ylim(0.1,1)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```

In [106]: import keras
from keras.applications.vgg16 import VGG16
from keras.preprocessing import image
from keras.applications.vgg16 import preprocess_input
from keras.models import Model
from keras.layers import Input, Dense, GlobalAveragePooling2D, Conv2D, MaxPooling2D, BatchNormalization, Dropout
import numpy as np

img_input = Input((32,32,3))
# Block 1
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv1')(img_input)
x = BatchNormalization()(x)
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv2')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block1_pool')(x)
x = Dropout(0.2)(x)

# Block 2
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv1')(x)
x = BatchNormalization()(x)
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv2')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block2_pool')(x)
x = Dropout(0.3)(x)

# Block 3
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv1')(x)
x = BatchNormalization()(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv2')(x)
x = BatchNormalization()(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv3')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block3_pool')(x)
x = Dropout(0.4)(x)
'''

# Block 4
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block4_pool')(x)

# Block 5
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_

```

```
conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_
conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)
'''

# Final Block
x = GlobalAveragePooling2D()(x)
x = Dense(32, activation='relu')(x)
x = Dropout(0.4)(x)
predictions = Dense(10, activation='softmax')(x)

model = Model(inputs=img_input, outputs=predictions)
model.summary()
```

Layer (type)	Output Shape	Param #
input_41 (InputLayer)	(None, 32, 32, 3)	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
batch_normalization_133 (Batch Normalization)	(None, 32, 32, 64)	256
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
batch_normalization_134 (Batch Normalization)	(None, 32, 32, 64)	256
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
dropout_28 (Dropout)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
batch_normalization_135 (Batch Normalization)	(None, 16, 16, 128)	512
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
batch_normalization_136 (Batch Normalization)	(None, 16, 16, 128)	512
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
dropout_29 (Dropout)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
batch_normalization_137 (Batch Normalization)	(None, 8, 8, 256)	1024
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
batch_normalization_138 (Batch Normalization)	(None, 8, 8, 256)	1024
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
batch_normalization_139 (Batch Normalization)	(None, 8, 8, 256)	1024
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
dropout_30 (Dropout)	(None, 4, 4, 256)	0
global_average_pooling2d_28 (Global Average Pooling)	(None, 256)	0
dense_101 (Dense)	(None, 32)	8224
dropout_31 (Dropout)	(None, 32)	0
dense_102 (Dense)	(None, 10)	330
Total params: 1,748,650		
Trainable params: 1,746,346		
Non-trainable params: 2,304		



```
In [103]: BATCH_SIZE = 128
MAX_EPOCH = 50
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.SGD(lr=1e-3),
              metrics=['accuracy'])
history_SGD_4 = model.fit(x_train_VGG, y_train_VGG,
                        batch_size=BATCH_SIZE,
                        epochs=MAX_EPOCH,
                        verbose=1,
                        validation_data=(x_test_VGG, y_test_VGG))
score = model.evaluate(x_test_VGG, y_test_VGG, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```



Train on 50000 samples, validate on 10000 samples

Epoch 1/50  
50000/50000 [=====] - 15s 301us/step - loss:  
2.2862 - acc: 0.1914 - val\_loss: 1.9612 - val\_acc: 0.2939

Epoch 2/50  
50000/50000 [=====] - 12s 241us/step - loss:  
2.0321 - acc: 0.2555 - val\_loss: 1.8533 - val\_acc: 0.3342

Epoch 3/50  
50000/50000 [=====] - 12s 241us/step - loss:  
1.9277 - acc: 0.2953 - val\_loss: 1.7633 - val\_acc: 0.3548

Epoch 4/50  
50000/50000 [=====] - 12s 238us/step - loss:  
1.8424 - acc: 0.3305 - val\_loss: 1.8035 - val\_acc: 0.3575

Epoch 5/50  
50000/50000 [=====] - 12s 242us/step - loss:  
1.7782 - acc: 0.3575 - val\_loss: 1.6776 - val\_acc: 0.3914

Epoch 6/50  
50000/50000 [=====] - 12s 239us/step - loss:  
1.7208 - acc: 0.3780 - val\_loss: 1.6545 - val\_acc: 0.4093

Epoch 7/50  
50000/50000 [=====] - 12s 241us/step - loss:  
1.6820 - acc: 0.3936 - val\_loss: 1.5162 - val\_acc: 0.4487

Epoch 8/50  
50000/50000 [=====] - 12s 241us/step - loss:  
1.6386 - acc: 0.4118 - val\_loss: 1.5732 - val\_acc: 0.4513

Epoch 9/50  
50000/50000 [=====] - 12s 242us/step - loss:  
1.6075 - acc: 0.4229 - val\_loss: 1.5855 - val\_acc: 0.4465

Epoch 10/50  
50000/50000 [=====] - 12s 240us/step - loss:  
1.5779 - acc: 0.4347 - val\_loss: 1.6130 - val\_acc: 0.4485

Epoch 11/50  
50000/50000 [=====] - 12s 238us/step - loss:  
1.5516 - acc: 0.4454 - val\_loss: 1.6987 - val\_acc: 0.4287

Epoch 12/50  
50000/50000 [=====] - 12s 247us/step - loss:  
1.5231 - acc: 0.4556 - val\_loss: 1.5415 - val\_acc: 0.4600

Epoch 13/50  
50000/50000 [=====] - 12s 243us/step - loss:  
1.5020 - acc: 0.4643 - val\_loss: 1.5993 - val\_acc: 0.4606

Epoch 14/50  
50000/50000 [=====] - 12s 245us/step - loss:  
1.4826 - acc: 0.4696 - val\_loss: 1.3991 - val\_acc: 0.5136

Epoch 15/50  
50000/50000 [=====] - 12s 241us/step - loss:  
1.4623 - acc: 0.4765 - val\_loss: 1.5623 - val\_acc: 0.4711

Epoch 16/50  
50000/50000 [=====] - 12s 245us/step - loss:  
1.4412 - acc: 0.4876 - val\_loss: 1.4687 - val\_acc: 0.4897

Epoch 17/50  
50000/50000 [=====] - 12s 242us/step - loss:  
1.4177 - acc: 0.4970 - val\_loss: 1.5310 - val\_acc: 0.4811

Epoch 18/50  
50000/50000 [=====] - 12s 241us/step - loss:  
1.4064 - acc: 0.4969 - val\_loss: 1.3754 - val\_acc: 0.5208



Epoch 19/50  
50000/50000 [=====] - 12s 240us/step - loss:

```
1.3845 - acc: 0.5085 - val_loss: 1.4238 - val_acc: 0.4955
Epoch 20/50
50000/50000 [=====] - 12s 241us/step - loss:
1.3724 - acc: 0.5144 - val_loss: 1.8246 - val_acc: 0.4044
Epoch 21/50
50000/50000 [=====] - 12s 246us/step - loss:
1.3544 - acc: 0.5183 - val_loss: 1.3383 - val_acc: 0.5432
Epoch 22/50
50000/50000 [=====] - 12s 246us/step - loss:
1.3365 - acc: 0.5257 - val_loss: 1.3913 - val_acc: 0.5201
Epoch 23/50
50000/50000 [=====] - 12s 243us/step - loss:
1.3299 - acc: 0.5300 - val_loss: 1.3439 - val_acc: 0.5356
Epoch 24/50
50000/50000 [=====] - 12s 240us/step - loss:
1.3156 - acc: 0.5353 - val_loss: 1.2246 - val_acc: 0.5657
Epoch 25/50
50000/50000 [=====] - 12s 247us/step - loss:
1.3045 - acc: 0.5362 - val_loss: 1.4800 - val_acc: 0.5014
Epoch 26/50
50000/50000 [=====] - 12s 239us/step - loss:
1.2868 - acc: 0.5449 - val_loss: 1.3365 - val_acc: 0.5379
Epoch 27/50
50000/50000 [=====] - 12s 239us/step - loss:
1.2759 - acc: 0.5515 - val_loss: 1.3220 - val_acc: 0.5447
Epoch 28/50
50000/50000 [=====] - 12s 242us/step - loss:
1.2697 - acc: 0.5504 - val_loss: 1.1870 - val_acc: 0.5920
Epoch 29/50
50000/50000 [=====] - 12s 243us/step - loss:
1.2563 - acc: 0.5570 - val_loss: 1.2085 - val_acc: 0.5820
Epoch 30/50
50000/50000 [=====] - 12s 241us/step - loss:
1.2392 - acc: 0.5650 - val_loss: 1.2537 - val_acc: 0.5582
Epoch 31/50
50000/50000 [=====] - 12s 239us/step - loss:
1.2289 - acc: 0.5646 - val_loss: 1.4619 - val_acc: 0.5319
Epoch 32/50
50000/50000 [=====] - 12s 243us/step - loss:
1.2186 - acc: 0.5718 - val_loss: 1.3392 - val_acc: 0.5353
Epoch 33/50
50000/50000 [=====] - 12s 241us/step - loss:
1.2058 - acc: 0.5741 - val_loss: 1.2020 - val_acc: 0.5802
Epoch 34/50
50000/50000 [=====] - 12s 241us/step - loss:
1.2034 - acc: 0.5755 - val_loss: 1.2782 - val_acc: 0.5603
Epoch 35/50
50000/50000 [=====] - 12s 240us/step - loss:
1.1898 - acc: 0.5808 - val_loss: 1.1882 - val_acc: 0.5885
Epoch 36/50
50000/50000 [=====] - 12s 243us/step - loss:
1.1809 - acc: 0.5855 - val_loss: 1.4110 - val_acc: 0.5323
Epoch 37/50
50000/50000 [=====] - 12s 244us/step - loss:
1.1726 - acc: 0.5872 - val_loss: 1.1987 - val_acc: 0.5755
Epoch 38/50
50000/50000 [=====] - 12s 238us/step - loss:
```

```
1.1600 - acc: 0.5933 - val_loss: 1.1558 - val_acc: 0.6027
Epoch 39/50
50000/50000 [=====] - 12s 243us/step - loss:
1.1542 - acc: 0.5948 - val_loss: 1.2833 - val_acc: 0.5739
Epoch 40/50
50000/50000 [=====] - 12s 238us/step - loss:
1.1450 - acc: 0.5959 - val_loss: 1.1337 - val_acc: 0.6033
Epoch 41/50
50000/50000 [=====] - 12s 247us/step - loss:
1.1324 - acc: 0.6039 - val_loss: 1.0621 - val_acc: 0.6264
Epoch 42/50
50000/50000 [=====] - 12s 245us/step - loss:
1.1246 - acc: 0.6063 - val_loss: 1.2674 - val_acc: 0.5613
Epoch 43/50
50000/50000 [=====] - 12s 246us/step - loss:
1.1163 - acc: 0.6092 - val_loss: 1.0566 - val_acc: 0.6312
Epoch 44/50
50000/50000 [=====] - 12s 241us/step - loss:
1.1160 - acc: 0.6072 - val_loss: 1.3641 - val_acc: 0.5547
Epoch 45/50
50000/50000 [=====] - 12s 244us/step - loss:
1.1031 - acc: 0.6141 - val_loss: 1.0987 - val_acc: 0.6205
Epoch 46/50
50000/50000 [=====] - 12s 245us/step - loss:
1.0957 - acc: 0.6160 - val_loss: 1.1913 - val_acc: 0.5870
Epoch 47/50
50000/50000 [=====] - 12s 245us/step - loss:
1.0886 - acc: 0.6167 - val_loss: 1.2157 - val_acc: 0.5844
Epoch 48/50
50000/50000 [=====] - 12s 243us/step - loss:
1.0789 - acc: 0.6196 - val_loss: 1.1787 - val_acc: 0.5887
Epoch 49/50
50000/50000 [=====] - 12s 242us/step - loss:
1.0727 - acc: 0.6228 - val_loss: 1.1375 - val_acc: 0.6022
Epoch 50/50
50000/50000 [=====] - 12s 246us/step - loss:
1.0658 - acc: 0.6236 - val_loss: 1.3320 - val_acc: 0.5622
Test loss: 1.3320011768341065
Test accuracy: 0.5622
```

```
In [172]: # summarize history for loss
plt.plot(history_SGD_4.history['loss'])
plt.plot(history_SGD_4.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.ylim(0,2.2)
plt.legend(['train', 'test'], loc='upper left')
plt.show()

# summarize history for accuracy
plt.plot(history_SGD_4.history['acc'])
plt.plot(history_SGD_4.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.ylim(0.1,1)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
In [98]: BATCH_SIZE = 128
MAX_EPOCH = 50
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adagrad(),
              metrics=['accuracy'])
history_Adagrad = model.fit(x_train_VGG, y_train_VGG,
                           batch_size=BATCH_SIZE,
                           epochs=MAX_EPOCH,
                           verbose=1,
                           validation_data=(x_test_VGG, y_test_VGG))
score = model.evaluate(x_test_VGG, y_test_VGG, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 50000 samples, validate on 10000 samples

Epoch 1/50  
50000/50000 [=====] - 14s 288us/step - loss: 1.5609 - acc: 0.4375 - val\_loss: 1.7541 - val\_acc: 0.3955

Epoch 2/50  
50000/50000 [=====] - 12s 238us/step - loss: 1.1060 - acc: 0.6158 - val\_loss: 1.0400 - val\_acc: 0.6258

Epoch 3/50  
50000/50000 [=====] - 12s 238us/step - loss: 0.9401 - acc: 0.6768 - val\_loss: 0.9982 - val\_acc: 0.6581

Epoch 4/50  
50000/50000 [=====] - 12s 237us/step - loss: 0.8251 - acc: 0.7177 - val\_loss: 0.8195 - val\_acc: 0.7123

Epoch 5/50  
50000/50000 [=====] - 12s 238us/step - loss: 0.7415 - acc: 0.7489 - val\_loss: 0.6885 - val\_acc: 0.7575

Epoch 6/50  
50000/50000 [=====] - 12s 238us/step - loss: 0.6700 - acc: 0.7727 - val\_loss: 0.6448 - val\_acc: 0.7715

Epoch 7/50  
50000/50000 [=====] - 12s 240us/step - loss: 0.6114 - acc: 0.7975 - val\_loss: 0.6118 - val\_acc: 0.7984

Epoch 8/50  
50000/50000 [=====] - 12s 242us/step - loss: 0.5661 - acc: 0.8126 - val\_loss: 0.6019 - val\_acc: 0.7983

Epoch 9/50  
50000/50000 [=====] - 12s 239us/step - loss: 0.5244 - acc: 0.8262 - val\_loss: 0.6264 - val\_acc: 0.7978

Epoch 10/50  
50000/50000 [=====] - 12s 239us/step - loss: 0.4912 - acc: 0.8374 - val\_loss: 0.5505 - val\_acc: 0.8239

Epoch 11/50  
50000/50000 [=====] - 12s 241us/step - loss: 0.4571 - acc: 0.8479 - val\_loss: 0.5229 - val\_acc: 0.8268

Epoch 12/50  
50000/50000 [=====] - 12s 242us/step - loss: 0.4224 - acc: 0.8610 - val\_loss: 0.5471 - val\_acc: 0.8268

Epoch 13/50  
50000/50000 [=====] - 12s 239us/step - loss: 0.4002 - acc: 0.8677 - val\_loss: 0.5382 - val\_acc: 0.8299

Epoch 14/50  
50000/50000 [=====] - 12s 242us/step - loss: 0.3726 - acc: 0.8752 - val\_loss: 0.5386 - val\_acc: 0.8367

Epoch 15/50  
50000/50000 [=====] - 12s 239us/step - loss: 0.3540 - acc: 0.8825 - val\_loss: 0.5038 - val\_acc: 0.8461

Epoch 16/50  
50000/50000 [=====] - 12s 243us/step - loss: 0.3337 - acc: 0.8905 - val\_loss: 0.5370 - val\_acc: 0.8383

Epoch 17/50  
50000/50000 [=====] - 12s 240us/step - loss: 0.3181 - acc: 0.8932 - val\_loss: 0.5678 - val\_acc: 0.8365

Epoch 18/50  
50000/50000 [=====] - 12s 242us/step - loss: 0.2957 - acc: 0.9029 - val\_loss: 0.5513 - val\_acc: 0.8437


Epoch 19/50  
50000/50000 [=====] - 12s 242us/step - loss:

```
0.2817 - acc: 0.9067 - val_loss: 0.5216 - val_acc: 0.8553
Epoch 20/50
50000/50000 [=====] - 12s 240us/step - loss:
0.2665 - acc: 0.9117 - val_loss: 0.5434 - val_acc: 0.8485
Epoch 21/50
50000/50000 [=====] - 12s 241us/step - loss:
0.2487 - acc: 0.9180 - val_loss: 0.5597 - val_acc: 0.8471
Epoch 22/50
50000/50000 [=====] - 12s 238us/step - loss:
0.2399 - acc: 0.9210 - val_loss: 0.5528 - val_acc: 0.8492
Epoch 23/50
50000/50000 [=====] - 12s 243us/step - loss:
0.2279 - acc: 0.9252 - val_loss: 0.5891 - val_acc: 0.8426
Epoch 24/50
50000/50000 [=====] - 12s 237us/step - loss:
0.2146 - acc: 0.9293 - val_loss: 0.5728 - val_acc: 0.8503
Epoch 25/50
50000/50000 [=====] - 12s 235us/step - loss:
0.2086 - acc: 0.9312 - val_loss: 0.5611 - val_acc: 0.8544
Epoch 26/50
50000/50000 [=====] - 12s 238us/step - loss:
0.1922 - acc: 0.9374 - val_loss: 0.5833 - val_acc: 0.8546
Epoch 27/50
50000/50000 [=====] - 12s 238us/step - loss:
0.1822 - acc: 0.9405 - val_loss: 0.5801 - val_acc: 0.8539
Epoch 28/50
50000/50000 [=====] - 12s 239us/step - loss:
0.1780 - acc: 0.9431 - val_loss: 0.5786 - val_acc: 0.8541
Epoch 29/50
50000/50000 [=====] - 12s 238us/step - loss:
0.1668 - acc: 0.9445 - val_loss: 0.5734 - val_acc: 0.8570
Epoch 30/50
50000/50000 [=====] - 12s 240us/step - loss:
0.1559 - acc: 0.9486 - val_loss: 0.5968 - val_acc: 0.8545
Epoch 31/50
50000/50000 [=====] - 12s 243us/step - loss:
0.1481 - acc: 0.9518 - val_loss: 0.6054 - val_acc: 0.8521
Epoch 32/50
50000/50000 [=====] - 12s 238us/step - loss:
0.1427 - acc: 0.9545 - val_loss: 0.5759 - val_acc: 0.8612
Epoch 33/50
50000/50000 [=====] - 12s 239us/step - loss:
0.1391 - acc: 0.9546 - val_loss: 0.6258 - val_acc: 0.8574
Epoch 34/50
50000/50000 [=====] - 12s 239us/step - loss:
0.1335 - acc: 0.9571 - val_loss: 0.6089 - val_acc: 0.8561
Epoch 35/50
50000/50000 [=====] - 12s 240us/step - loss:
0.1271 - acc: 0.9594 - val_loss: 0.6101 - val_acc: 0.8570
Epoch 36/50
50000/50000 [=====] - 12s 240us/step - loss:
0.1205 - acc: 0.9617 - val_loss: 0.6217 - val_acc: 0.8629
Epoch 37/50
50000/50000 [=====] - 12s 239us/step - loss:
0.1135 - acc: 0.9639 - val_loss: 0.6323 - val_acc: 0.8580
Epoch 38/50
50000/50000 [=====] - 12s 239us/step - loss:
```

```
0.1143 - acc: 0.9635 - val_loss: 0.6320 - val_acc: 0.8618
Epoch 39/50
50000/50000 [=====] - 12s 240us/step - loss:
0.1126 - acc: 0.9651 - val_loss: 0.6266 - val_acc: 0.8585
Epoch 40/50
50000/50000 [=====] - 12s 237us/step - loss:
0.1056 - acc: 0.9666 - val_loss: 0.6324 - val_acc: 0.8594
Epoch 41/50
50000/50000 [=====] - 12s 238us/step - loss:
0.0928 - acc: 0.9698 - val_loss: 0.6656 - val_acc: 0.8607
Epoch 42/50
50000/50000 [=====] - 12s 237us/step - loss:
0.0945 - acc: 0.9700 - val_loss: 0.6556 - val_acc: 0.8597
Epoch 43/50
50000/50000 [=====] - 12s 238us/step - loss:
0.0951 - acc: 0.9705 - val_loss: 0.6680 - val_acc: 0.8607
Epoch 44/50
50000/50000 [=====] - 12s 240us/step - loss:
0.0914 - acc: 0.9712 - val_loss: 0.6884 - val_acc: 0.8572
Epoch 45/50
50000/50000 [=====] - 12s 240us/step - loss:
0.0861 - acc: 0.9733 - val_loss: 0.6704 - val_acc: 0.8625
Epoch 46/50
50000/50000 [=====] - 12s 236us/step - loss:
0.0861 - acc: 0.9736 - val_loss: 0.6689 - val_acc: 0.8632
Epoch 47/50
50000/50000 [=====] - 12s 238us/step - loss:
0.0778 - acc: 0.9753 - val_loss: 0.6811 - val_acc: 0.8582
Epoch 48/50
50000/50000 [=====] - 12s 240us/step - loss:
0.0768 - acc: 0.9760 - val_loss: 0.6894 - val_acc: 0.8634
Epoch 49/50
50000/50000 [=====] - 12s 240us/step - loss:
0.0756 - acc: 0.9764 - val_loss: 0.6727 - val_acc: 0.8653
Epoch 50/50
50000/50000 [=====] - 12s 239us/step - loss:
0.0719 - acc: 0.9779 - val_loss: 0.6950 - val_acc: 0.8580
Test loss: 0.6949959478616714
Test accuracy: 0.858
```



```
In [173]: # summarize history for loss
plt.plot(history_Adagrad.history['loss'])
plt.plot(history_Adagrad.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.ylim(0,2.2)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
# summarize history for accuracy
plt.plot(history_Adagrad.history['acc'])
plt.plot(history_Adagrad.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.ylim(0.1,1)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
In [105]: BATCH_SIZE = 128
MAX_EPOCH = 50
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Nadam(),
              metrics=['accuracy'])
history_Nadam = model.fit(x_train_VGG, y_train_VGG,
                          batch_size=BATCH_SIZE,
                          epochs=MAX_EPOCH,
                          verbose=1,
                          validation_data=(x_test_VGG, y_test_VGG))
score = model.evaluate(x_test_VGG, y_test_VGG, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 50000 samples, validate on 10000 samples

Epoch 1/50  
50000/50000 [=====] - 16s 322us/step - loss: 1.5625 - acc: 0.4508 - val\_loss: 1.9672 - val\_acc: 0.4389

Epoch 2/50  
50000/50000 [=====] - 12s 249us/step - loss: 1.0304 - acc: 0.6474 - val\_loss: 1.2605 - val\_acc: 0.5755

Epoch 3/50  
50000/50000 [=====] - 13s 253us/step - loss: 0.8213 - acc: 0.7260 - val\_loss: 0.8337 - val\_acc: 0.7186

Epoch 4/50  
50000/50000 [=====] - 13s 254us/step - loss: 0.6931 - acc: 0.7735 - val\_loss: 0.7237 - val\_acc: 0.7626

Epoch 5/50  
50000/50000 [=====] - 13s 256us/step - loss: 0.6096 - acc: 0.8024 - val\_loss: 0.6445 - val\_acc: 0.7918

Epoch 6/50  
50000/50000 [=====] - 13s 251us/step - loss: 0.5340 - acc: 0.8278 - val\_loss: 0.5692 - val\_acc: 0.8096

Epoch 7/50  
50000/50000 [=====] - 13s 254us/step - loss: 0.4824 - acc: 0.8454 - val\_loss: 0.7835 - val\_acc: 0.7410

Epoch 8/50  
50000/50000 [=====] - 12s 249us/step - loss: 0.4305 - acc: 0.8631 - val\_loss: 0.5519 - val\_acc: 0.8297

Epoch 9/50  
50000/50000 [=====] - 13s 254us/step - loss: 0.3902 - acc: 0.8726 - val\_loss: 0.6172 - val\_acc: 0.8205

Epoch 10/50  
50000/50000 [=====] - 13s 253us/step - loss: 0.3593 - acc: 0.8830 - val\_loss: 0.5819 - val\_acc: 0.8310

Epoch 11/50  
50000/50000 [=====] - 13s 251us/step - loss: 0.3242 - acc: 0.8955 - val\_loss: 0.5308 - val\_acc: 0.8452

Epoch 12/50  
50000/50000 [=====] - 13s 254us/step - loss: 0.3042 - acc: 0.9016 - val\_loss: 0.5435 - val\_acc: 0.8447

Epoch 13/50  
50000/50000 [=====] - 13s 251us/step - loss: 0.2739 - acc: 0.9128 - val\_loss: 0.5698 - val\_acc: 0.8419

Epoch 14/50  
50000/50000 [=====] - 13s 254us/step - loss: 0.2558 - acc: 0.9168 - val\_loss: 0.5450 - val\_acc: 0.8475

Epoch 15/50  
50000/50000 [=====] - 12s 249us/step - loss: 0.2322 - acc: 0.9249 - val\_loss: 0.5153 - val\_acc: 0.8585

Epoch 16/50  
50000/50000 [=====] - 13s 253us/step - loss: 0.2181 - acc: 0.9305 - val\_loss: 0.5596 - val\_acc: 0.8500

Epoch 17/50  
50000/50000 [=====] - 13s 252us/step - loss: 0.2080 - acc: 0.9330 - val\_loss: 0.6785 - val\_acc: 0.8322



Epoch 18/50  
50000/50000 [=====] - 12s 249us/step - loss: 0.1945 - acc: 0.9375 - val\_loss: 0.5489 - val\_acc: 0.8504

Epoch 19/50  
50000/50000 [=====] - 13s 254us/step - loss:

```
0.1771 - acc: 0.9426 - val_loss: 0.7583 - val_acc: 0.8288
Epoch 20/50
50000/50000 [=====] - 13s 252us/step - loss:
0.1701 - acc: 0.9465 - val_loss: 0.6427 - val_acc: 0.8476
Epoch 21/50
50000/50000 [=====] - 13s 255us/step - loss:
0.1578 - acc: 0.9507 - val_loss: 0.6095 - val_acc: 0.8542
Epoch 22/50
50000/50000 [=====] - 13s 255us/step - loss:
0.1551 - acc: 0.9513 - val_loss: 0.6053 - val_acc: 0.8517
Epoch 23/50
50000/50000 [=====] - 13s 254us/step - loss:
0.1478 - acc: 0.9530 - val_loss: 0.6065 - val_acc: 0.8621
Epoch 24/50
50000/50000 [=====] - 13s 252us/step - loss:
0.1367 - acc: 0.9570 - val_loss: 0.6644 - val_acc: 0.8571
Epoch 25/50
50000/50000 [=====] - 13s 251us/step - loss:
0.1373 - acc: 0.9565 - val_loss: 0.5943 - val_acc: 0.8636
Epoch 26/50
50000/50000 [=====] - 13s 254us/step - loss:
0.1159 - acc: 0.9642 - val_loss: 0.6615 - val_acc: 0.8552
Epoch 27/50
50000/50000 [=====] - 12s 247us/step - loss:
0.1199 - acc: 0.9620 - val_loss: 0.6411 - val_acc: 0.8590
Epoch 28/50
50000/50000 [=====] - 12s 246us/step - loss:
0.1165 - acc: 0.9633 - val_loss: 0.6589 - val_acc: 0.8561
Epoch 29/50
50000/50000 [=====] - 12s 245us/step - loss:
0.1133 - acc: 0.9650 - val_loss: 0.6851 - val_acc: 0.8496
Epoch 30/50
50000/50000 [=====] - 13s 251us/step - loss:
0.1101 - acc: 0.9660 - val_loss: 0.7552 - val_acc: 0.8379
Epoch 31/50
50000/50000 [=====] - 12s 248us/step - loss:
0.0992 - acc: 0.9691 - val_loss: 0.6487 - val_acc: 0.8642
Epoch 32/50
50000/50000 [=====] - 13s 253us/step - loss:
0.1064 - acc: 0.9664 - val_loss: 0.6558 - val_acc: 0.8654
Epoch 33/50
50000/50000 [=====] - 12s 248us/step - loss:
0.0931 - acc: 0.9706 - val_loss: 0.7756 - val_acc: 0.8471
Epoch 34/50
50000/50000 [=====] - 12s 249us/step - loss:
0.0957 - acc: 0.9698 - val_loss: 0.8121 - val_acc: 0.8513
Epoch 35/50
50000/50000 [=====] - 12s 249us/step - loss:
0.0959 - acc: 0.9704 - val_loss: 0.6664 - val_acc: 0.8580
Epoch 36/50
50000/50000 [=====] - 13s 250us/step - loss:
0.0888 - acc: 0.9730 - val_loss: 0.8507 - val_acc: 0.8360
Epoch 37/50
50000/50000 [=====] - 12s 247us/step - loss:
0.0827 - acc: 0.9747 - val_loss: 0.7951 - val_acc: 0.8380
Epoch 38/50
50000/50000 [=====] - 12s 247us/step - loss:
```

```
0.0889 - acc: 0.9719 - val_loss: 0.7148 - val_acc: 0.8517
Epoch 39/50
50000/50000 [=====] - 12s 241us/step - loss:
0.0905 - acc: 0.9718 - val_loss: 0.6577 - val_acc: 0.8651
Epoch 40/50
50000/50000 [=====] - 12s 236us/step - loss:
0.0767 - acc: 0.9748 - val_loss: 0.7139 - val_acc: 0.8634
Epoch 41/50
50000/50000 [=====] - 12s 249us/step - loss:
0.0786 - acc: 0.9754 - val_loss: 0.7351 - val_acc: 0.8618
Epoch 42/50
50000/50000 [=====] - 13s 251us/step - loss:
0.0776 - acc: 0.9760 - val_loss: 0.7491 - val_acc: 0.8643
Epoch 43/50
50000/50000 [=====] - 13s 252us/step - loss:
0.0792 - acc: 0.9756 - val_loss: 0.7160 - val_acc: 0.8537
Epoch 44/50
50000/50000 [=====] - 13s 251us/step - loss:
0.0751 - acc: 0.9766 - val_loss: 0.7801 - val_acc: 0.8607
Epoch 45/50
50000/50000 [=====] - 13s 252us/step - loss:
0.0719 - acc: 0.9776 - val_loss: 0.6928 - val_acc: 0.8673
Epoch 46/50
50000/50000 [=====] - 12s 248us/step - loss:
0.0712 - acc: 0.9777 - val_loss: 0.6921 - val_acc: 0.8669
Epoch 47/50
50000/50000 [=====] - 12s 248us/step - loss:
0.0738 - acc: 0.9773 - val_loss: 0.6879 - val_acc: 0.8681
Epoch 48/50
50000/50000 [=====] - 12s 249us/step - loss:
0.0689 - acc: 0.9779 - val_loss: 0.7538 - val_acc: 0.8630
Epoch 49/50
50000/50000 [=====] - 12s 248us/step - loss:
0.0637 - acc: 0.9800 - val_loss: 0.7102 - val_acc: 0.8673
Epoch 50/50
50000/50000 [=====] - 12s 248us/step - loss:
0.0663 - acc: 0.9794 - val_loss: 0.6751 - val_acc: 0.8749
Test loss: 0.6751188353247941
Test accuracy: 0.8749
```

```
In [174]: # summarize history for loss
plt.plot(history_Nadam.history['loss'])
plt.plot(history_Nadam.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.ylim(0,2.2)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
# summarize history for accuracy
plt.plot(history_Nadam.history['acc'])
plt.plot(history_Nadam.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.ylim(0.1,1)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
In [107]: BATCH_SIZE = 128
MAX_EPOCH = 50
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.RMSprop(),
              metrics=['accuracy'])
history_RMSprop = model.fit(x_train_VGG, y_train_VGG,
                           batch_size=BATCH_SIZE,
                           epochs=MAX_EPOCH,
                           verbose=1,
                           validation_data=(x_test_VGG, y_test_VGG))
score = model.evaluate(x_test_VGG, y_test_VGG, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
Train on 50000 samples, validate on 10000 samples
Epoch 1/50
50000/50000 [=====] - 16s 329us/step - loss:
1.5991 - acc: 0.4330 - val_loss: 1.6797 - val_acc: 0.4591
Epoch 2/50
50000/50000 [=====] - 12s 245us/step - loss:
1.0683 - acc: 0.6324 - val_loss: 1.0910 - val_acc: 0.6369
Epoch 3/50
50000/50000 [=====] - 12s 238us/step - loss:
0.8587 - acc: 0.7119 - val_loss: 0.9857 - val_acc: 0.6677
Epoch 4/50
50000/50000 [=====] - 12s 241us/step - loss:
0.7303 - acc: 0.7604 - val_loss: 0.7986 - val_acc: 0.7431
Epoch 5/50
50000/50000 [=====] - 12s 239us/step - loss:
0.6381 - acc: 0.7930 - val_loss: 0.7648 - val_acc: 0.7548
Epoch 6/50
50000/50000 [=====] - 12s 241us/step - loss:
0.5716 - acc: 0.8161 - val_loss: 0.6257 - val_acc: 0.7989
Epoch 7/50
50000/50000 [=====] - 12s 238us/step - loss:
0.5098 - acc: 0.8353 - val_loss: 0.6098 - val_acc: 0.8039
Epoch 8/50
50000/50000 [=====] - 12s 240us/step - loss:
0.4670 - acc: 0.8488 - val_loss: 0.5470 - val_acc: 0.8191
Epoch 9/50
50000/50000 [=====] - 12s 243us/step - loss:
0.4266 - acc: 0.8628 - val_loss: 0.5741 - val_acc: 0.8242
Epoch 10/50
50000/50000 [=====] - 12s 247us/step - loss:
0.3908 - acc: 0.8749 - val_loss: 0.6792 - val_acc: 0.7996
Epoch 11/50
50000/50000 [=====] - 12s 249us/step - loss:
0.3563 - acc: 0.8867 - val_loss: 0.7010 - val_acc: 0.8056
Epoch 12/50
50000/50000 [=====] - 13s 252us/step - loss:
0.3306 - acc: 0.8944 - val_loss: 0.6557 - val_acc: 0.8216
Epoch 13/50
50000/50000 [=====] - 12s 248us/step - loss:
0.3031 - acc: 0.9023 - val_loss: 0.6659 - val_acc: 0.8140
Epoch 14/50
50000/50000 [=====] - 12s 243us/step - loss:
0.2829 - acc: 0.9088 - val_loss: 0.6305 - val_acc: 0.8356
Epoch 15/50
50000/50000 [=====] - 12s 247us/step - loss:
0.2636 - acc: 0.9149 - val_loss: 0.7954 - val_acc: 0.7886
Epoch 16/50
50000/50000 [=====] - 12s 248us/step - loss:
0.2476 - acc: 0.9224 - val_loss: 0.7604 - val_acc: 0.8160
Epoch 17/50
50000/50000 [=====] - 12s 248us/step - loss:
0.2309 - acc: 0.9254 - val_loss: 0.5690 - val_acc: 0.8586
Epoch 18/50
50000/50000 [=====] - 12s 246us/step - loss:
0.2167 - acc: 0.9315 - val_loss: 0.6179 - val_acc: 0.8488
Epoch 19/50
50000/50000 [=====] - 12s 248us/step - loss:
```



0.2015 - acc: 0.9362 - val\_loss: 0.6492 - val\_acc: 0.8466  
Epoch 20/50  
50000/50000 [=====] - 12s 247us/step - loss:  
0.1989 - acc: 0.9383 - val\_loss: 0.7450 - val\_acc: 0.8202  
Epoch 21/50  
50000/50000 [=====] - 12s 247us/step - loss:  
0.1876 - acc: 0.9416 - val\_loss: 0.6410 - val\_acc: 0.8440  
Epoch 22/50  
50000/50000 [=====] - 12s 249us/step - loss:  
0.1725 - acc: 0.9467 - val\_loss: 0.6248 - val\_acc: 0.8521  
Epoch 23/50  
50000/50000 [=====] - 12s 248us/step - loss:  
0.1636 - acc: 0.9493 - val\_loss: 0.7388 - val\_acc: 0.8244  
Epoch 24/50  
50000/50000 [=====] - 12s 249us/step - loss:  
0.1571 - acc: 0.9511 - val\_loss: 0.7895 - val\_acc: 0.8358  
Epoch 25/50  
50000/50000 [=====] - 12s 245us/step - loss:  
0.1533 - acc: 0.9543 - val\_loss: 0.7159 - val\_acc: 0.8476  
Epoch 26/50  
50000/50000 [=====] - 12s 248us/step - loss:  
0.1446 - acc: 0.9566 - val\_loss: 0.7043 - val\_acc: 0.8384  
Epoch 27/50  
50000/50000 [=====] - 12s 246us/step - loss:  
0.1375 - acc: 0.9582 - val\_loss: 0.6664 - val\_acc: 0.8552  
Epoch 28/50  
50000/50000 [=====] - 12s 245us/step - loss:  
0.1335 - acc: 0.9597 - val\_loss: 0.6555 - val\_acc: 0.8522  
Epoch 29/50  
50000/50000 [=====] - 12s 249us/step - loss:  
0.1286 - acc: 0.9607 - val\_loss: 0.6874 - val\_acc: 0.8434  
Epoch 30/50  
50000/50000 [=====] - 12s 247us/step - loss:  
0.1251 - acc: 0.9623 - val\_loss: 0.8071 - val\_acc: 0.8376  
Epoch 31/50  
50000/50000 [=====] - 12s 246us/step - loss:  
0.1198 - acc: 0.9648 - val\_loss: 0.6501 - val\_acc: 0.8632  
Epoch 32/50  
50000/50000 [=====] - 12s 247us/step - loss:  
0.1179 - acc: 0.9658 - val\_loss: 0.7089 - val\_acc: 0.8425  
Epoch 33/50  
50000/50000 [=====] - 12s 246us/step - loss:  
0.1129 - acc: 0.9656 - val\_loss: 0.8513 - val\_acc: 0.8277  
Epoch 34/50  
50000/50000 [=====] - 12s 247us/step - loss:  
0.1117 - acc: 0.9664 - val\_loss: 0.9202 - val\_acc: 0.8273  
Epoch 35/50  
50000/50000 [=====] - 12s 242us/step - loss:  
0.1104 - acc: 0.9679 - val\_loss: 0.6386 - val\_acc: 0.8636  
Epoch 36/50  
50000/50000 [=====] - 12s 246us/step - loss:  
0.1033 - acc: 0.9694 - val\_loss: 0.7566 - val\_acc: 0.8524  
Epoch 37/50  
50000/50000 [=====] - 12s 249us/step - loss:  
0.1005 - acc: 0.9713 - val\_loss: 0.7331 - val\_acc: 0.8556  
Epoch 38/50  
50000/50000 [=====] - 12s 245us/step - loss:

```
0.1003 - acc: 0.9713 - val_loss: 0.7536 - val_acc: 0.8563
Epoch 39/50
50000/50000 [=====] - 12s 249us/step - loss:
0.0994 - acc: 0.9720 - val_loss: 0.7282 - val_acc: 0.8556
Epoch 40/50
50000/50000 [=====] - 12s 248us/step - loss:
0.0982 - acc: 0.9712 - val_loss: 0.7373 - val_acc: 0.8599
Epoch 41/50
50000/50000 [=====] - 12s 247us/step - loss:
0.0936 - acc: 0.9731 - val_loss: 0.7248 - val_acc: 0.8597
Epoch 42/50
50000/50000 [=====] - 12s 245us/step - loss:
0.0882 - acc: 0.9742 - val_loss: 0.7132 - val_acc: 0.8597
Epoch 43/50
50000/50000 [=====] - 12s 243us/step - loss:
0.0904 - acc: 0.9737 - val_loss: 0.7408 - val_acc: 0.8550
Epoch 44/50
50000/50000 [=====] - 12s 246us/step - loss:
0.0857 - acc: 0.9748 - val_loss: 0.8124 - val_acc: 0.8380
Epoch 45/50
50000/50000 [=====] - 12s 243us/step - loss:
0.0914 - acc: 0.9738 - val_loss: 0.7211 - val_acc: 0.8587
Epoch 46/50
50000/50000 [=====] - 12s 247us/step - loss:
0.0830 - acc: 0.9759 - val_loss: 0.8645 - val_acc: 0.8535
Epoch 47/50
50000/50000 [=====] - 12s 250us/step - loss:
0.0845 - acc: 0.9761 - val_loss: 0.6940 - val_acc: 0.8693
Epoch 48/50
50000/50000 [=====] - 12s 245us/step - loss:
0.0858 - acc: 0.9757 - val_loss: 0.7605 - val_acc: 0.8647
Epoch 49/50
50000/50000 [=====] - 12s 242us/step - loss:
0.0846 - acc: 0.9764 - val_loss: 0.8083 - val_acc: 0.8552
Epoch 50/50
50000/50000 [=====] - 12s 248us/step - loss:
0.0771 - acc: 0.9781 - val_loss: 0.8476 - val_acc: 0.8480
Test loss: 0.8475979769527913
Test accuracy: 0.848
```

```
In [175]: # summarize history for loss
plt.plot(history_RMSprop.history['loss'])
plt.plot(history_RMSprop.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.ylim(0,2.2)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
# summarize history for accuracy
plt.plot(history_RMSprop.history['acc'])
plt.plot(history_RMSprop.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.ylim(0.1,1)
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

